Realized Volatility Forecasting with Time Series Foundation Models

Alessio Brini*

March 12, 2026

Abstract

Realized volatility forecasting is central to risk management and portfolio allocation, yet no study has systematically assessed whether pretrained time series foundation models can match established econometric benchmarks for this task. We evaluate four zero-shot foundation models against six econometric specifications on 50 assets from the VOLARE dataset (equities, FX, futures; 2015–2026) at horizons of 1, 5, and 22 days, using Diebold–Mariano tests and Model Confidence Sets. The best foundation models match or exceed the best econometric benchmarks, but rankings depend on loss function, horizon, and asset class. At the daily horizon, Log-HAR retains a narrow QLIKE edge while foundation models dominate on MSE and out-of-sample R^2; at longer horizons, Moirai 2.0 leads across all metrics. Performance is heterogeneous across architectures: not all foundation models are competitive, and rankings shift across asset classes.

Keywords: Realized volatility, foundation models, zero-shot forecasting, HAR, model confidence set **JEL Codes:** C22, C53, C55, G17

*Duke University. Email: `alessio.brini@duke.edu`.

1 Introduction

Forecasting realized volatility is central to risk management, derivative pricing, and portfolio allocation. Since the work of Andersen and Bollerslev (1998) and Barndorff-Nielsen and Shephard (2002), realized volatility constructed from high-frequency returns has become the standard model-free measure of ex-post price variation. The Heterogeneous Autoregressive (HAR) model of Corsi (2009) is the dominant forecasting benchmark: its three-component structure, aggregating past realized volatility at daily, weekly, and monthly frequencies, provides a parsimonious approximation to the long-memory dynamics that characterize volatility. Extensions incorporating jump variation (Andersen et al., 2007), semivariance asymmetries (Patton and Sheppard, 2015), and measurement error corrections (Bollerslev et al., 2016) have yielded mixed practical gains (Clements and Preve, 2021; Branco et al., 2024). Machine learning methods, including neural networks (Bucci, 2020; Brini and Toscano, 2025), random forests (Luong and Dokuchaev, 2018), graph-based approaches (Zhang et al., 2024, 2025), and convolutional architectures (Moreno-Pino and Zohren, 2024), have also been applied, but a broader assessment by Christensen et al. (2023) concludes that machine learning offers incremental improvements rather than a paradigm shift.

A different class of models has emerged over the past two years: time series foundation models (TSFMs). These are large pretrained transformers (Vaswani et al., 2017) trained on billions of time series observations across domains, capable of producing forecasts for previously unseen series without any retraining, a capability known as zero-shot forecasting. Leading examples include Chronos (Ansari et al., 2024), Moirai (Woo et al., 2024b), Lag-Llama (Rasul et al., 2024), and Kronos (Woo et al., 2024a), several of which have released second-generation versions with improved architectures (Ansari et al., 2025; Liu et al., 2025a). Surveys document the rapid pace of development in this area (Liang et al., 2024; Ye et al., 2024; Miller et al., 2024). Outside finance, Carriero et al. (2024) applied zero-shot language models to macroeconomic forecasting and found that they do not outperform traditional econometric baselines, raising the question of whether similar conclusions hold for financial

time series.

Applications of TSFMs to financial time series remain limited. Goel et al. (2025) tested TimesFM on realized volatility for 20 stocks and found that fine-tuning was necessary for the model to compete with HAR; zero-shot performance was poor. Rahimikia et al. (2025) evaluated several TSFMs on daily stock returns and reported uniformly negative zero-shot results. However, realized volatility differs from returns in ways that may favor foundation models: it is strictly positive, exhibits well-documented mean reversion and long-memory dependence (Andersen et al., 2003), and follows a distribution closer to the macroeconomic and physical series on which TSFMs were pretrained. To our knowledge, no study has yet evaluated multiple foundation models on realized volatility with formal statistical testing.

This paper conducts the first systematic comparison of zero-shot TSFMs against established econometric benchmarks for realized volatility forecasting. We evaluate four TSFMs against six econometric specifications across 50 assets spanning three asset classes (equities, FX, futures) and three forecast horizons ($h = 1, 5, 22$ days), using the recently released VOLARE dataset (Cipollini et al., 2026). The zero-shot setting, with no domain-specific training, isolates the value of general time series pretraining. We apply Diebold–Mariano tests and Model Confidence Set analysis (Hansen et al., 2011) to control for multiple comparisons.

Zero-shot TSFMs are competitive with the best econometric benchmarks, contrasting with prior negative findings on daily returns (Rahimikia et al., 2025) and on an earlier-generation TSFM (Goel et al., 2025). Model rankings depend on the loss function and horizon: Log-HAR retains a narrow QLIKE advantage at $h = 1$ for equities, but TSFMs dominate on MSE, MAE, and out-of-sample R^2 at all horizons, and Moirai 2.0 achieves the lowest QLIKE at $h = 5$ and $h = 22$. Model Confidence Set analysis shows that Moirai 2.0 enters the 90% MCS for every equity at $h = 5$ and for 97.5% at $h = 22$, while Log-HAR enters for 90% at $h = 1$. Rankings are not stable across asset classes, with different TSFMs dominating for equities, FX, and futures. Levels-based HAR extensions produce near-zero

forecasts that inflate QLIKE, a practical limitation that log-space and TSFM approaches avoid.

The remainder of this paper is organized as follows. Section 2 reviews the related literature. Section 3 describes the VOLARE dataset. Section 4 presents the econometric and foundation model specifications, along with the forecast evaluation framework. Section 5 reports the empirical results. Section 6 contains robustness checks, and Section 7 concludes.

2 Literature Review

2.1 Realized Volatility Modeling

The theory of realized volatility originates with Andersen and Bollerslev (1998) and Andersen et al. (2001), who showed that the sum of squared intraday returns provides a consistent, nonparametric estimator of the integrated variance of asset prices. Barndorff-Nielsen and Shephard (2002) formalized this result in a semimartingale framework, establishing that realized variance converges to integrated variance as the sampling frequency increases. Subsequent work identified the key stylized facts of realized volatility: approximate log-normality, long-memory dependence with a fractional integration parameter $d \approx 0.4$, and slow mean reversion (Andersen et al., 2003).

At ultra-high sampling frequencies, microstructure noise from bid-ask bounce and discrete price changes biases the realized variance estimator upward. Barndorff-Nielsen et al. (2008) introduced the realized kernel as a noise-consistent alternative that remains valid in the presence of market microstructure effects.

The HAR model of Corsi (2009) became the standard forecasting benchmark for realized volatility. By including daily, weekly, and monthly RV components as regressors, the model approximates long-memory behavior through a cascade of time scales corresponding to heterogeneous trading horizons. The specification is parsimonious (three regressors plus a constant) yet achieves accuracy comparable to far more complex alternatives.

A large body of work has extended the HAR framework. Andersen et al. (2007) and Barndorff-Nielsen and Shephard (2004) separated realized variance into continuous and jump components using bipower variation, yielding the HAR-J model. Patton and Sheppard (2015) and Barndorff-Nielsen et al. (2010) decomposed realized variance into positive and negative semivariance to capture asymmetric responses to upside and downside moves. Bollerslev et al. (2016) introduced the HARQ model, which interacts the daily RV regressor with realized quarticity to account for time-varying measurement error. Clements and Preve (2021) provides a practical guide to implementing and comparing these extensions.

Long-memory models offer a complementary approach. ARFIMA models (Baillie et al., 1996; Izzeldin et al., 2019) explicitly parameterize the fractional integration order and remain competitive at longer forecast horizons where the slow mean reversion of realized volatility becomes the dominant dynamic.

Machine learning methods have been applied to realized volatility forecasting with mixed results. Bucci (2020) reports that feedforward and recurrent neural networks outperform HAR on a sample of U.S. stocks. Luong and Dokuchaev (2018) finds gains from random forests. Moreno-Pino and Zohren (2024) introduces DeepVol, a convolutional neural network that processes raw order book data. In contrast, Branco et al. (2024) concludes that simple linear models are difficult to beat after correcting for multiple testing, and Christensen et al. (2023) finds that machine learning offers only incremental improvements over standard econometric methods. Recent surveys by Gunnarsson et al. (2024) and Leushuis and Petkov (2026) cover this literature.

2.2 Time Series Foundation Models

Time series foundation models are large pretrained models that produce forecasts for arbitrary time series without task-specific training, analogous to large language models for text (Liang et al., 2024). These models are trained on vast corpora of time series drawn from diverse domains (weather, energy, retail, transport, macroeconomics) and learn general

temporal patterns that transfer to unseen series at inference time.

Three architectural families dominate the first generation of TSFMs. Chronos (Ansari et al., 2024) converts continuous values into discrete tokens via quantile binning and uses a T5 encoder-decoder architecture (Vaswani et al., 2017) for autoregressive forecasting; T5 ("Text-to-Text Transfer Transformer") is a sequence-to-sequence model that reads the full input through an encoder and generates output tokens autoregressively through a decoder. Moirai (Woo et al., 2024b) handles arbitrary numbers of variates through an "any-variate" attention mechanism and uses mixture distribution outputs for uncertainty quantification. Lag-Llama (Rasul et al., 2024) adapts the LLaMA decoder-only architecture for probabilistic time series forecasting using lag-based tokenization. Kronos (Woo et al., 2024a) employs kernel-based tokenization, representing each continuous value as a weighted combination of learned kernel basis functions, and processes the resulting embeddings through a T5 backbone.

Second-generation models appeared in 2025. Ansari et al. (2025) introduced Chronos-2 and Chronos-Bolt: the former adds multivariate support and covariate handling, while Chronos-Bolt is a distilled variant, meaning its parameters were trained to approximate the full Chronos-2 model's behavior at a fraction of the computational cost, producing direct quantile forecasts in a single forward pass and running up to 250 times faster than the original. Moirai 2.0 (Liu et al., 2025a) demonstrated that smaller, better-trained models can match or exceed their larger predecessors using multi-token prediction and improved tokenization. Other models include MOMENT (Goswami et al., 2024), a masked-encoder architecture designed for multiple time series tasks (forecasting, classification, anomaly detection), Timer (Liu et al., 2024), and Sundial (Liu et al., 2025b), each exploring different design choices for temporal representation and output distribution.

Standardized benchmarks have been developed to evaluate TSFMs across domains. GIFT-Eval (Aksu et al., 2024) provides a unified evaluation protocol across multiple datasets and forecast horizons. FEV-Bench (Shchur et al., 2025) focuses on forecasting evaluation with

emphasis on calibration. TSFM-Bench (Li et al., 2024) compares models across zero-shot and fine-tuned settings. A consistent finding is that no single model dominates across all domains and horizons, which motivates domain-specific evaluations such as the one we undertake here. Adler et al. (2025) raises additional concerns about the calibration of TSFM prediction intervals, finding systematic miscoverage across several models.

2.3 Foundation Models for Volatility Forecasting

The application of TSFMs to volatility forecasting is recent and limited. The closest prior work is Goel et al. (2025), who tested TimesFM 1.0 on realized volatility for 20 stocks. They found that zero-shot TimesFM underperformed HAR and that fine-tuning was necessary to achieve competitive accuracy. Our paper differs in several respects: we evaluate four foundation models from different families (Chronos-Bolt, Moirai 2.0, Lag-Llama, and Kronos), use newer model generations, and apply Model Confidence Set analysis to control for multiple comparisons. The same research group also applied TimesFM to Value-at-Risk forecasting (Goel et al., 2024).

Rahimikia et al. (2025) evaluated Chronos, TimesFM, Moirai, and MOMENT on daily stock returns and found that zero-shot TSFMs consistently underperformed ARIMA. However, daily returns differ from realized volatility in fundamental ways: returns are approximately symmetric, near-white-noise, and lack the strong serial dependence that characterizes realized volatility. Realized volatility is strictly positive, mean-reverting, and displays long-memory dynamics, properties that bring it closer to the macroeconomic and physical time series on which TSFMs were pretrained. This distinction suggests that negative results on returns need not carry over to realized measures.

Other financial applications include stock price forecasting (Łaniewski and Ślepaczuk, 2025; Valeyre and Aboura, 2024), foreign exchange volatility modeling (Nguyen et al., 2025), and multivariate financial time series (Marconi, 2025). Finance-specific foundation models have also been proposed, including FinCast (Zhu et al., 2025) and a finance-oriented model

also named Kronos (Shi et al., 2025). The Kronos-T5 model we use in our evaluation (Woo et al., 2024a) is the general-purpose variant released by Salesforce, distinct from the finance-specific model of Shi et al. (2025).

The question of whether to fine-tune or apply TSFMs zero-shot is itself an active area of research. Das et al. (2025) explored in-context fine-tuning for TimesFM, while Germán-Morales et al. (2024) and Gupta et al. (2024) investigated parameter-efficient adaptation via Low-Rank Adaptation (LoRA; Hu et al., 2022). Testing zero-shot performance first establishes an important baseline: if a TSFM already performs well without adaptation, the marginal value of fine-tuning may be limited, and the zero-shot workflow offers practical advantages in simplicity and computational cost.

We contribute to this emerging literature by providing the first evaluation that compares multiple time series foundation models against a broad set of econometric benchmarks for realized volatility, with formal statistical testing via Diebold–Mariano tests (Diebold and Mariano, 1995) and Model Confidence Sets (Hansen et al., 2011).

3 The VOLARE Dataset

We use the VOLARE (VOLatility Archive for Realized Estimates) dataset of Cipollini et al. (2026), which provides daily realized volatility measures for a broad cross-section of financial assets. VOLARE is constructed from ultra-high-frequency tick data sourced from Kibot, covering 40 U.S. equities, 5 major currency pairs, and 5 commodity and index futures contracts. The equity sample begins on January 2, 2015 and runs through January 30, 2026, yielding 2,786 trading days per stock. The FX and futures samples begin on September 25, 2009 (4,242 and 4,228 trading days, respectively).

For each asset-day, VOLARE provides realized measures computed at multiple sampling frequencies. All measures used in this paper are computed at the 5-minute sampling frequency: realized variance (RV), bipower variation (BPV), positive and negative realized

semivariance (RS^+, RS^-), and realized quarticity (RQ). All values are expressed in decimal squared returns; for instance, a typical daily RV for a U.S. equity is approximately 1.9×10^{-4}, corresponding to an annualized volatility of roughly 22% ($\sqrt{252 \times 1.9 \times 10^{-4}} \approx 0.22$). VOLARE is well suited for our multi-asset comparison because it constructs all realized measures using a uniform methodology across assets, ensuring that cross-asset differences in model rankings reflect genuine forecasting performance rather than measurement inconsistencies.

Our equity sample comprises the 40 stocks in VOLARE that span the full 2015–2026 period: AAPL, ADBE, AMD, AMGN, AMZN, AXP, BA, CAT, CRM, CSCO, CVX, DIS, GE, GOOGL, GS, HD, HON, IBM, JNJ, JPM, KO, MCD, META, MMM, MRK, MSFT, NFLX, NKE, NVDA, ORCL, PG, PM, SHW, TRV, TSLA, UNH, V, VZ, WMT, and XOM. These cover all 11 Global Industry Classification Standard (GICS) sectors. The FX sample consists of five major currency pairs (AUDUSD, EURUSD, GBPUSD, USDCAD, USDJPY) and the futures sample covers five contracts (Corn, Crude Oil, E-mini S&P 500, Gold, Natural Gas).

Table 1 reports descriptive statistics for the 5-minute realized variance across the three asset classes. As expected for a strictly positive quantity, realized variance is right-skewed across all asset classes, with heavy tails (kurtosis ranging from 35 to over 3,000) confirming that extreme volatility episodes are a persistent feature of the data. The first-order autocorrelation of daily RV, denoted ρ_1, averages 0.606 for equities, consistent with the well-documented persistence of volatility; the 22-day autocorrelation (ρ_{22}, i.e., the sample autocorrelation at lag 22) averages 0.085, indicating slow decay. The FX pairs display RV approximately two orders of magnitude smaller than equities (cross-sectional average $\bar{RV} \approx 3 \times 10^{-5}$ vs. 2×10^{-4}), with comparable persistence (cross-sectional average $\bar{\rho}_1 = 0.52$, where the bar denotes averaging across assets). Futures exhibit the widest cross-asset heterogeneity: the E-mini S&P 500 (ES) has $\rho_1 = 0.779$, while Gold (GC) shows near-zero autocorrelation ($\rho_1 = -0.001$), presenting a natural stress test for forecasting models.

Table 1: Descriptive statistics for daily realized variance measured at the 5-minute frequency. The equity panel reports cross-sectional averages across 40 stocks; FX and futures report individual assets. ρ_k denotes the sample autocorrelation at lag k (e.g., ρ_{22} is the autocorrelation at 22 trading days). All RV values are in decimal squared returns ($\times 10^{-4}$).

	N	Mean	Median	Skew	Kurt	ρ_1	ρ_{22}
Panel A: Equities (cross-sectional average, 40 stocks)							
Average	2,786	1.91	1.12	13.5	269	0.606	0.085
Panel B: Foreign Exchange (5 pairs)							
AUDUSD	4,241	0.48	0.37	11.3	232	0.577	0.138
EURUSD	4,241	0.28	0.22	5.0	51	0.570	0.263
GBPUSD	4,242	0.32	0.24	27.5	1,117	0.484	0.061
USDCAD	4,240	0.25	0.19	4.3	35	0.654	0.368
USDJPY	4,239	0.33	0.22	14.4	349	0.325	0.063
Panel C: Futures (5 contracts)							
Corn (C)	4,185	2.48	1.54	12.3	238	0.178	0.079
Crude Oil (CL)	4,223	6.87	3.87	57.6	3,546	0.195	0.047
E-mini S&P (ES)	4,224	1.11	0.53	12.9	212	0.779	0.110
Gold (GC)	4,206	3.63	2.32	62.3	3,963	−0.001	−0.000
Natural Gas (NG)	4,221	10.30	6.35	19.0	522	0.187	0.267

4 Methodology

4.1 Econometric Benchmarks

We consider six econometric specifications that span the main approaches to realized volatility forecasting: the HAR family and its extensions, which exploit multi-scale aggregation, and ARFIMA, which explicitly models long memory.

HAR Model. The Heterogeneous Autoregressive model of Corsi (2009) captures the multi-scale persistence of realized volatility by aggregating past RV at daily, weekly, and monthly frequencies:

$$RV_{t+h} = \beta_0 + \beta_1 RV_t + \beta_2 RV_{t-5:t} + \beta_3 RV_{t-22:t} + \varepsilon_{t+h}, \tag{1}$$

where $RV_{t-k:t} \equiv k^{-1} \sum_{i=0}^{k-1} RV_{t-i}$ denotes the average realized volatility over the previous k days. We use the original Corsi (2009) specification, where the weekly and monthly components include overlapping lags (i.e., the weekly component averages days t through $t-4$, not the non-overlapping "rotated" version of Patton and Sheppard 2015 that separates lags 2–5 from 6–22). This three-component structure approximates long-memory behavior through the actions of heterogeneous market participants operating at different time horizons (Corsi, 2009).

HAR-J Model. The HAR-J model augments the baseline HAR with a jump component to separate continuous and discontinuous variation (Andersen et al., 2007):

$$RV_{t+h} = \beta_0 + \beta_1 RV_t + \beta_2 RV_{t-5:t} + \beta_3 RV_{t-22:t} + \beta_J J_t + \varepsilon_{t+h}, \tag{2}$$

where $J_t = \max(RV_t - BPV_t, 0)$ measures the jump component as the positive part of the difference between realized variance and bipower variation (Barndorff-Nielsen and Shephard, 2004). A negative coefficient on J_t indicates that large jumps reduce, rather than increase, future volatility.

HAR-RS Model. The HAR-RS model decomposes realized variance into positive and negative semivariance components to capture asymmetric volatility responses (Patton and Sheppard, 2015):

$$RV_{t+h} = \beta_0 + \beta_1^+ RS_t^+ + \beta_1^- RS_t^- + \beta_2^+ RS_{t-5:t}^+ + \beta_2^- RS_{t-5:t}^- + \beta_3^+ RS_{t-22:t}^+ + \beta_3^- RS_{t-22:t}^- + \varepsilon_{t+h}, \tag{3}$$

where RS_t^+ and RS_t^- denote good and bad realized semivariance, respectively, and $RS_t^+ + RS_t^- = RV_t$ by construction (Barndorff-Nielsen et al., 2010). The six-regressor structure allows upside and downside risk to follow separate dynamics at each aggregation frequency.

HARQ Model. The HARQ model of Bollerslev et al. (2016) accounts for time-varying measurement error in realized volatility by interacting the daily RV regressor with the square root of realized quarticity:

$$RV_{t+h} = \beta_0 + (\beta_1 + \beta_1^Q \sqrt{RQ_t})\, RV_t + \beta_2 RV_{t-5:t} + \beta_3 RV_{t-22:t} + \varepsilon_{t+h}, \tag{4}$$

where RQ_t is the realized quarticity. The interaction term attenuates the daily RV signal when measurement noise is high, as indicated by large values of RQ_t.

Log-HAR Model. The Log-HAR model applies the HAR specification to log-transformed realized volatility (Andersen et al., 2003):

$$\log RV_{t+h} = \beta_0 + \beta_1 \log RV_t + \beta_2 \log RV_{t-5:t} + \beta_3 \log RV_{t-22:t} + \varepsilon_{t+h}. \tag{5}$$

The logarithmic transformation brings the RV distribution closer to normality, improving OLS efficiency. Point forecasts in levels are recovered via bias-corrected retransformation: $\widehat{RV}_{t+h} = \exp(\hat{y}_{t+h} + \hat{\sigma}^2/2)$, where $\hat{\sigma}^2$ is the estimated residual variance, following the standard log-normal adjustment of Andersen et al. (2003).

ARFIMA Model. The ARFIMA model captures the well-documented long-memory property of realized volatility through fractional differencing (Baillie et al., 1996):

$$(1-L)^d\, \Phi(L)\, (\log RV_t - \mu) = \Theta(L)\, \varepsilon_t, \tag{6}$$

where $d \in (0, 0.5)$ is the fractional integration parameter, $\Phi(L)$ and $\Theta(L)$ are autoregressive and moving average lag polynomials of orders p and q (with $p, q \leq 2$ selected by AIC), and L is the lag operator. The parameter d is estimated by maximum likelihood. The fractional differencing filter is truncated at $K = 100$ lags.

All HAR-family models are estimated by OLS with Newey–West heteroskedasticity and

autocorrelation consistent standard errors. The ARFIMA model is re-estimated every 22 trading days to balance computational cost against parameter staleness.

4.2 Time Series Foundation Models

We evaluate four time series foundation models (TSFMs) in a zero-shot setting, applied directly to realized volatility series without any domain-specific training or fine-tuning. The zero-shot approach is the most practically relevant deployment mode: it requires no labeled financial data, no model retraining, and can be applied immediately to any new asset. Among the four TSFMs we evaluate, Chronos-Bolt and Moirai 2.0 are inference-only architectures with no public fine-tuning API. Lag-Llama supports fine-tuning through its GluonTS estimator interface, and the full Chronos-2 model (distinct from the distilled Bolt variant) provides built-in LoRA fine-tuning (Hu et al., 2022) via its `fit()` method (Ansari et al., 2025); we leave these adaptation strategies for future work.[1]

General pipeline. Because TSFMs originate in the machine learning literature, we briefly describe the forecasting pipeline in terms familiar to econometricians. Every TSFM follows three steps: *tokenization*, *encoding*, and *decoding*. First, the raw time series, here a vector of 512 daily RV observations, is converted into a sequence of discrete or continuous *tokens*, the model's internal representation of the data. Tokenization serves the same role as variable construction in an econometric model: it determines which features of the history the model can "see." Different TSFMs tokenize differently (quantile binning, lag features, kernel embeddings), and we describe each below. Second, the token sequence is processed by a *transformer* (Vaswani et al., 2017), a neural network architecture in which each position in the sequence attends to all other positions via learned attention weights. The attention mechanism allows the model to capture dependencies at arbitrary lags without imposing

[1] We originally planned to include TimesFM 2.5 (Das et al., 2024), but its publicly released model weights are incompatible with the current version of the `timesfm` Python package (v1.3.0), producing dimension mismatches during inference that prevent reliable forecasting.

the specific functional form that HAR prescribes (daily, weekly, monthly components). The transformer's parameters are fixed during our evaluation; they were learned during pretraining on large external corpora of time series observations drawn from weather, energy, retail, transport, and macroeconomic domains. Third, the transformer's output is passed through a *decoding head* that maps the internal representation back to forecasts in the original scale. Decoding heads differ across models: some output quantiles of the predictive distribution directly, others parameterize an explicit probability distribution (e.g., Student-t) from which forecast trajectories are sampled. In both cases, we extract the median as the point forecast, analogous to reporting the conditional median from a quantile regression.

Chronos-Bolt. Chronos-Bolt is a distilled variant of the Chronos family (Ansari et al., 2024, 2025). The original Chronos model tokenizes continuous values by mapping them into 4,096 discrete bins via quantile binning, where each observation is replaced by the index of its nearest quantile in a reference distribution, much as one would discretize a continuous variable into histogram bins. The model then processes these token indices through a T5 encoder-decoder transformer and generates forecasts autoregressively, predicting one token at a time. Chronos-Bolt replaces this two-step architecture with an encoder-only transformer that reads the entire input context and produces all forecast quantiles in a single forward pass, without iterative decoding. This yields a 250× speedup over the original while maintaining competitive accuracy (Ansari et al., 2025). The decoding head outputs quantile forecasts directly (e.g., the 10th, 50th, and 90th percentiles), bypassing the need to sample from a token distribution. We use both the small (48M parameters) and base (205M parameters) checkpoints[2] and take the median (0.5 quantile) as the point forecast.

Moirai 2.0. Moirai 2.0 (Liu et al., 2025a) builds on the original Moirai architecture (Woo et al., 2024b). It tokenizes the input series by dividing it into fixed-length *patches* (contiguous blocks of consecutive observations, e.g., 32 days each) and projects each patch into the

[2] `amazon/chronos-bolt-small` and `amazon/chronos-bolt-base` on HuggingFace.

transformer's embedding space via a learned linear map. This is analogous to constructing non-overlapping moving-window summaries of the series, except that the summary is learned rather than hand-specified. The transformer is decoder-only: it processes patches causally (each patch attends only to itself and earlier patches), preserving the temporal ordering. Moirai 2.0 uses *multi-token prediction*, generating multiple future patches in a single pass rather than one step at a time. The decoding head outputs nine quantile levels (0.05, 0.1, 0.2, 0.3, 0.5, 0.7, 0.8, 0.9, 0.95) per forecast step. We use the small checkpoint[3] and extract the median quantile as the point forecast.

Lag-Llama. Lag-Llama (Rasul et al., 2024) is a decoder-only foundation model based on the LLaMA architecture, adapted for univariate probabilistic time series forecasting. Unlike the patch-based models above, Lag-Llama tokenizes each time step individually using a set of *lag features*: at each position t, the input vector consists of the current value and a fixed set of lagged values $(y_{t-1}, y_{t-2}, y_{t-7}, y_{t-14}, \ldots)$, concatenated with time-related covariates. This lag-based design is conceptually closer to an autoregressive econometric specification, since the model explicitly receives lagged values as inputs, though the transformer learns the mapping from lags to forecasts nonparametrically rather than through a linear projection. The decoding head parameterizes a Student-t mixture distribution at each forecast step: it outputs the location, scale, and degrees-of-freedom parameters for each mixture component. Forecasts are generated by sampling trajectories from this distribution and taking the median across samples. We use the pretrained checkpoint (8 layers, 4 attention heads, 36-dimensional embeddings per head), sample 20 trajectories, and report the median as the point forecast.

Kronos. Kronos (Woo et al., 2024a) uses *kernel-based tokenization*: each continuous value is represented as a weighted combination of a set of learned kernel basis functions, producing a soft embedding rather than a hard bin assignment. The embedding is then processed by a T5 encoder-decoder transformer backbone. At decoding time, the model samples token-level

[3] `Salesforce/moirai-2.0-R-small` on HuggingFace.

outputs from the learned distribution and maps them back to continuous values through the inverse kernel transform. We use the tiny checkpoint[4] (approximately 8M parameters) with a sample count of 5 and take the median prediction as the point forecast.

Zero-shot protocol. All four TSFMs receive a rolling context window of 512 daily RV observations as input, with no additional features or covariates. The 512-day context length is the maximum supported by all four architectures' positional encodings and is the default recommended setting; it corresponds to approximately two years of trading data, providing enough history to capture the long-memory dynamics of realized volatility. By contrast, the 252-day rolling estimation window for econometric models follows the convention of using one year of daily data, which is standard in the HAR literature (Corsi, 2009; Clements and Preve, 2021). The one-year window balances the need for sufficient data to estimate HAR's four parameters reliably against the risk of including stale observations from a different volatility regime (Patton and Sheppard, 2015; Bollerslev et al., 2016). This reflects each approach's recommended configuration rather than an equal-information comparison; we view this as the practically relevant setup, since a practitioner would deploy each model with its standard settings. We assess sensitivity to the context window length in Section 6.

For multi-step horizons ($h > 1$), the point forecast is the mean of the h-step output vector. No parameters are updated during the out-of-sample evaluation period; all forecasts are purely zero-shot. We extract the median rather than the mean of each model's predictive distribution because the median is the optimal point forecast under MAE loss and is more robust to the right-skewed forecast distributions that arise in volatility applications; for symmetric distributions, the two coincide. We acknowledge that the median is not the optimal point forecast under QLIKE loss (the conditional mean is); however, in practice, the median is the standard output of probabilistic TSFMs and produces more stable forecasts than the sample mean when the predictive distribution is heavy-tailed. Table 4.2 reports wall-clock timing for all models: econometric specifications require < 0.01 s per forecast

[4] `Salesforce/kronos-t5-tiny` on HuggingFace.

step on CPU, while TSFM inference ranges from 0.014 s (Chronos-Bolt-Small on GPU) to approximately 5 s (Kronos on CPU).

Model	Parameters	Device	Sec/step (h=1)	Total (1 asset)
Econometric models				
HAR	4	CPU	0.003	21 s
HAR-J	5	CPU	0.003	21 s
HAR-RS	7	CPU	0.004	27 s
HARQ	5	CPU	0.004	24 s
Log-HAR	4	CPU	0.004	30 s
ARFIMA	$\leq$7	CPU	0.006	45 s
Time series foundation models				
Chronos-Bolt-S	48M	GPU	0.014	1.6 min
Chronos-Bolt-B	205M	GPU	0.025	2.8 min
Moirai-2.0-S	14M	GPU	0.019	2.2 min
Lag-Llama	8M	GPU	0.084	9.5 min
Kronos	8M	CPU*	$\sim$5	$\sim$9.5 hr

Computational cost per forecast step. All experiments were run on the Duke Computing Cluster (DCC) using SLURM job scheduling on the `gpu-common` partition. Econometric models run on CPU; TSFMs run on a single NVIDIA `[GPU_PLACEHOLDER]` GPU. Seconds per step is the wall-clock time to produce one forecast for one asset-day, averaged across the equity sample ($n \approx$ 2,274 out-of-sample steps). Total time is the approximate wall-clock for one asset across three horizons. ARFIMA is re-estimated every 22 steps. Lag-Llama timing includes 20 sample paths per step. Kronos runs on CPU because its custom codebase lacks GPU support in our environment; its high per-step cost reflects both CPU execution and the OHLCV tokenization overhead. *CPU timing on cluster node (Intel Xeon, 4 cores).

4.3 Forecast Evaluation

Walk-forward design. We employ a walk-forward evaluation scheme with a rolling origin. Econometric models use a fixed estimation window of 252 trading days (approximately one year): at each origin, the model is estimated on the most recent 252 observations, and the window slides forward one day at a time, producing daily re-estimated forecasts. Foundation models follow the same walk-forward timing but use a 512-day context window and require no estimation step; only the context window slides forward. Because the TSFM context window is longer, the first TSFM forecast is produced 260 days later than the first econometric forecast; all model comparisons are conducted on the common out-of-sample period where

both econometric and TSFM forecasts are available. For the equity sample (2,786 days), this produces approximately 2,148 common out-of-sample forecasts per model–asset combination at $h = 1$; the FX and futures samples (4,240+ days) yield approximately 3,602.

Forecast horizons. We evaluate three forecast horizons: $h = 1$ (one day), $h = 5$ (one week), and $h = 22$ (one month). The forecast target at horizon h is the average realized volatility over the next h days, $h^{-1} \sum_{i=1}^{h} RV_{t+i}$. All models produce direct h-step forecasts. For HAR-family models, this means estimating a separate regression for each horizon, where the dependent variable is the h-day average of future RV; the model does not iterate one-step-ahead predictions forward. ARFIMA similarly produces direct multi-step forecasts from its estimated ARMA representation on log RV. For TSFMs, each model generates a full forecast trajectory of length h in a single inference call; the point forecast is the mean of this h-step output vector. This direct approach avoids the compounding of errors that arises in iterated multi-step forecasting.

Loss functions and statistical testing. Our primary loss function is the quasi-likelihood (QLIKE) loss of Patton (2011):

$$L_{\text{QLIKE}}(\hat{\sigma}_t^2, RV_t) = \frac{RV_t}{\hat{\sigma}_t^2} - \log \frac{RV_t}{\hat{\sigma}_t^2} - 1, \tag{7}$$

where $\hat{\sigma}_t^2$ is the forecast and RV_t is the realized value. Patton (2011) shows that QLIKE preserves the ranking of forecasts regardless of the choice of volatility proxy, making it the preferred metric when the true conditional variance is unobserved. We also report mean squared error (MSE), mean absolute error (MAE), and the out-of-sample R^2 defined as $R^2_{\text{OOS}} = 1 - \sum_t (RV_t - \hat{f}_t)^2 / \sum_t (RV_t - \bar{RV})^2$, where $\bar{RV}$ is the historical mean of realized variance over the estimation window. Unlike the in-sample R^2, R^2_{OOS} can take negative values, indicating that the model's forecasts are less accurate than a simple historical mean forecast. All forecasts are floored at 10^{-6} to prevent numerical instability in the QLIKE

calculation when models produce near-zero predictions.

Pairwise forecast comparisons use the Diebold and Mariano (1995) test with Newey–West variance estimation, applied to QLIKE loss differentials. To address the multiple comparison problem inherent in evaluating eleven models, we compute the Model Confidence Set (MCS) of Hansen et al. (2011), which identifies the subset of models whose forecasting ability cannot be statistically distinguished from the best model at a given significance level. We use the $T_{\max}$ test statistic with a block bootstrap (block length = 22 days, $B = 10{,}000$ replications) and significance level $\alpha = 0.10$.

We assess forecast efficiency using the Mincer–Zarnowitz regression (Mincer and Zarnowitz, 1969):

$$RV_t = \alpha + \beta \, \hat{f}_t + \varepsilon_t, \tag{8}$$

where $\hat{f}_t$ is the model's forecast. Under forecast optimality, $\alpha = 0$ and $\beta = 1$, meaning the forecast is unbiased and captures the correct scale of variation. We test the joint null $H_0 \colon \alpha = 0, \beta = 1$ using a Wald test with Newey–West standard errors and report cross-asset averages and rejection rates at the 5% level.

To examine whether relative forecast performance is stable over time, we apply the Giacomini and Rossi (2010) fluctuation test. For each model paired against the HAR benchmark, we compute a rolling DM statistic over a window of size $m = \lfloor 0.3 \times T \rfloor$, producing a time path of relative QLIKE performance. The test statistic is $\sup_t |S_t|$, compared against critical values from the distribution of the supremum of a standardized Brownian bridge. Rejection of the null indicates that the two models' relative performance changes significantly over time.

5 Empirical Results

This section presents the empirical results in four parts. We first examine aggregate forecast accuracy across loss functions and horizons, then analyze cross-asset and cross-market

heterogeneity. We next assess statistical significance through Diebold–Mariano tests, Model Confidence Sets, Mincer–Zarnowitz regressions, and Giacomini–Rossi fluctuation tests. Finally, we extend the analysis to realized covariance forecasting and global minimum variance portfolio construction.

5.1 Aggregate Forecast Accuracy

This subsection reports the overall forecasting performance averaged across assets within each asset class, without examining individual asset-level variation. The goal is to establish the broad ranking of models by loss function and horizon before turning to cross-asset heterogeneity and statistical testing in subsequent subsections.

Table 5.1 reports the cross-sectional average of per-asset loss functions across 40 equities. Figure 1 illustrates the forecasts of three representative models against realized values for four assets spanning equities and FX. The results reveal a systematic divergence between QLIKE and the other metrics.

Model	MSE ($\times 10^{-6}$)	MAE ($\times 10^{-4}$)	QLIKE	R^2_{OOS}
Panel: $h = 1$				
HAR	0.244	1.43	1.38†	0.335
HAR-J	0.242	1.42	1.16†	0.349
HAR-RS	0.265	1.52	3.73†	0.268
HARQ	0.297	1.49	16.46†	0.111
Log-HAR	0.257	1.31	**0.220**	0.238
ARFIMA	0.295	1.63	0.314	0.153
Chronos-Bolt-S	0.216	1.18	0.229	0.408
Chronos-Bolt-B	0.213	**1.17**	0.230	0.420
Moirai-2.0-S	**0.210**	1.19	0.226	**0.426**
Lag-Llama	0.291	1.65	0.402	0.127
Kronos	0.231	1.31	1.09†	0.348
Panel: $h = 5$				
HAR	0.163	1.40	2.70†	0.209
HAR-J	0.159	1.39	2.47†	0.231
HAR-RS	0.184	1.50	7.04†	0.101
HARQ	0.192	1.43	16.16†	0.024
Log-HAR	0.164	1.28	0.231	0.135
ARFIMA	0.170	1.44	0.244	0.178
Chronos-Bolt-S	0.127	1.09	0.213	0.355
Chronos-Bolt-B	0.128	1.09	0.216	0.346
Moirai-2.0-S	**0.122**	**1.07**	**0.188**	**0.384**
Lag-Llama	0.181	1.43	0.305	0.044
Kronos	0.143	1.22	0.867	0.243
Panel: $h = 22$				
HAR	0.117	1.47	6.06†	-0.031
HAR-J	0.117	1.46	5.92†	-0.015
HAR-RS	0.130	1.57	9.40†	-0.145
HARQ	0.131	1.48	12.42†	-0.184
Log-HAR	0.115	1.35	0.350	-0.044
ARFIMA	0.119	1.35	0.275	-0.014
Chronos-Bolt-S	0.100	1.13	0.322	0.085
Chronos-Bolt-B	0.101	1.13	0.325	0.077
Moirai-2.0-S	**0.098**	**1.07**	**0.257**	**0.110**
Lag-Llama	0.109	1.25	0.283	0.009
Kronos	0.108	1.25	0.913	0.012

Forecast accuracy for 40 U.S. equities (VOLARE). Values are cross-sectional averages of per-asset loss functions. Bold indicates the best value in each column within each panel. † marks QLIKE > 1: for HAR-RS and HARQ, this reflects near-zero forecasts from levels-based OLS; for Kronos, it reflects systematic forecast bias.

At the one-day horizon ($h = 1$), Log-HAR achieves the lowest average QLIKE (0.220),

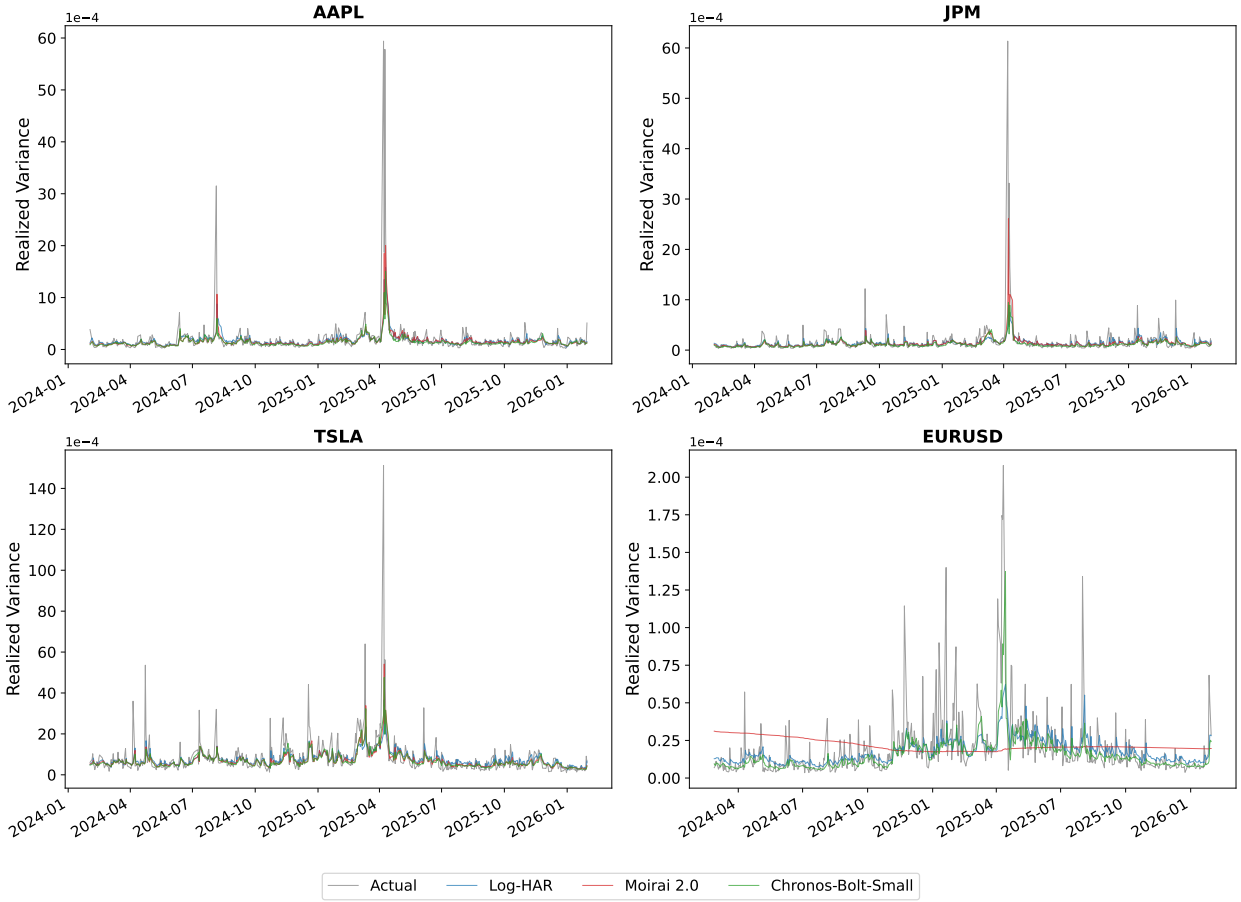

Figure 1: Forecast vs. actual realized variance at $h = 1$ for four representative assets. The last 500 out-of-sample observations are shown. Gray: actual RV; blue: Log-HAR; red: Moirai 2.0; green: Chronos-Bolt-Small.

followed closely by Moirai 2.0 (0.226) and Chronos-Bolt-Small (0.229). The gap between Log-HAR and the best TSFM is small (roughly 3% in relative terms) and is driven by the log-space formulation of Log-HAR, which guarantees positive forecasts and penalizes proportional forecast errors in a way that aligns with QLIKE's structure. On MSE, MAE, and R^2_{OOS}, however, the TSFMs dominate: Moirai 2.0 achieves the lowest MSE (0.210×10^{-6}), and the highest R^2_{OOS} (0.426), compared to 0.238 for Log-HAR. Chronos-Bolt-Base achieves the lowest MAE (1.17×10^{-4}). The out-of-sample R^2 gap between TSFMs and the best econometric model is 8–19 percentage points at $h = 1$, a substantial margin.

At the five-day horizon ($h = 5$), the TSFM advantage becomes more pronounced. Moirai 2.0 achieves the best score across all four metrics: MSE (0.122×10^{-6}), MAE

(1.07×10^{-4}), QLIKE (0.188), and R^2_{OOS} (0.384). The gap over Log-HAR in QLIKE terms is now 19%, compared to 3% at $h = 1$. Chronos-Bolt-Small achieves the second-lowest QLIKE (0.213), still below Log-HAR's 0.231. The econometric models that do best at $h = 1$ (HAR, HAR-J) deteriorate more rapidly at longer horizons, while the TSFMs maintain relatively stable accuracy.

At the monthly horizon ($h = 22$), Moirai 2.0 again dominates: QLIKE = 0.257, MSE = 0.098×10^{-6}, and $R^2_{\text{OOS}} = 0.110$. The econometric benchmarks now show near-zero or negative R^2_{OOS}, indicating that they barely outperform (or underperform) the historical mean at this horizon. Log-HAR achieves $R^2_{\text{OOS}} = -0.044$, while Chronos-Bolt-Small maintains a positive R^2_{OOS} of 0.085. This horizon-dependent pattern is consistent with the hypothesis that TSFMs' large context windows (512 days) capture long-range dependencies more effectively than the HAR's fixed three-component structure.

Levels-based HAR extensions are unstable in this evaluation, consistent with Clements and Preve (2021), who document that unconstrained OLS estimation of HAR variants can produce negative or near-zero forecasts. HAR-RS and HARQ produce QLIKE values of 3.73 and 16.46 at $h = 1$, respectively, rising to 9.40 and 12.42 at $h = 22$. These inflated values arise because the OLS-estimated models occasionally produce near-zero forecasts for RV_{t+h}, which are floored at 10^{-6} for QLIKE evaluation. Across the full sample, no forecast is literally negative, but the floor binds for a small fraction of observations. Since QLIKE = $RV/\hat{RV} - \log(RV/\hat{RV}) - 1$, a forecast near zero generates an extremely large loss whenever the realized value is far from zero, and a handful of such observations can dominate the cross-sectional average. Across the equity sample, the 10^{-6} floor binds for fewer than 0.1% of TSFM forecast–date pairs but for 2–5% of HAR-RS and HARQ pairs, explaining their inflated QLIKE. Table A reports cross-sectional medians, which are robust to these outliers; the qualitative ranking of models is unchanged. Log-HAR and ARFIMA avoid this problem entirely by operating in log space, where forecasts are positive by construction. TSFMs also produce strictly positive forecasts, since their training distributions are bounded below.

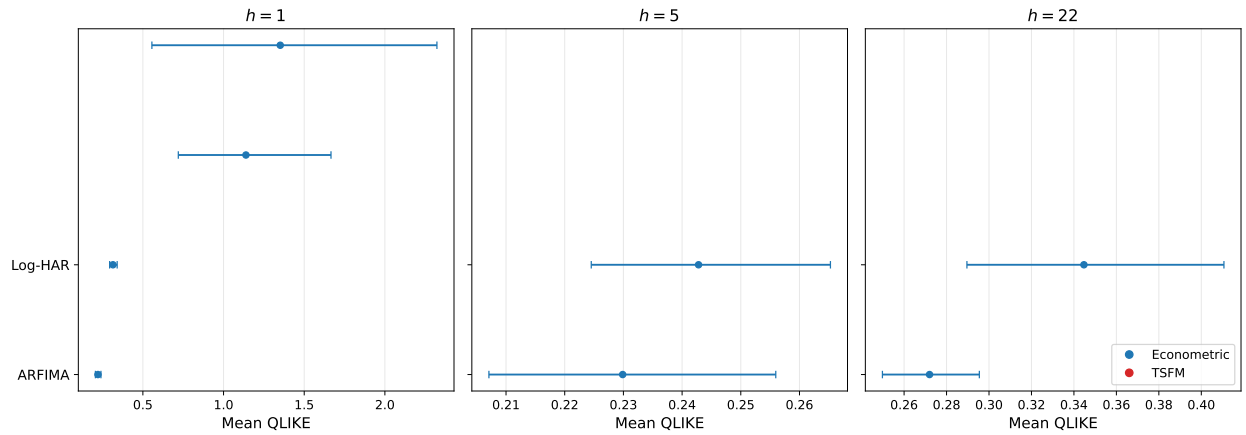

Figure 2: Bootstrap 95% confidence intervals for cross-sectional mean QLIKE across 40 equities. Blue: econometric models; red: TSFMs. Models with mean QLIKE > 2 are excluded for readability. Dots indicate point estimates; horizontal bars span the 2.5th to 97.5th percentile of the bootstrap distribution (10,000 draws).

This practical advantage may be as important as any differences in statistical accuracy: a model that occasionally produces degenerate forecasts is difficult to deploy in production risk systems.

Figure 2 visualizes the bootstrap confidence intervals for all models across the three horizons. Bootstrap confidence intervals (10,000 draws, percentile method) on the cross-sectional mean QLIKE confirm the statistical significance of the key rankings. At $h = 1$, the 95% CI for Log-HAR QLIKE is $[0.204, 0.239]$, overlapping with Moirai 2.0's $[0.209, 0.247]$ and Chronos-Bolt-Small's $[0.211, 0.251]$, consistent with the MCS finding that these models are statistically indistinguishable. At $h = 5$, Moirai 2.0's CI $[0.172, 0.209]$ lies entirely below Log-HAR's $[0.208, 0.257]$, confirming the TSFM advantage at this horizon. At $h = 22$, Moirai 2.0 $[0.238, 0.277]$ and Lag-Llama $[0.260, 0.306]$ both separate cleanly from Log-HAR $[0.294, 0.417]$.

5.2 Cross-Asset and Cross-Market Heterogeneity

Tables 5.2 and 5.2 extend the analysis to foreign exchange and futures markets, revealing cross-asset-class heterogeneity in model rankings.

Model	MSE ($\times 10^{-8}$)	MAE ($\times 10^{-4}$)	QLIKE	R^2_{OOS}
Panel: $h = 1$				
HAR	0.174	0.14	0.177	0.241
HAR-J	0.170	0.14	0.196	0.258
HAR-RS	0.192	0.15	0.249	0.180
HARQ	0.196	0.14	0.628	0.167
Log-HAR	0.174	0.13	**0.163**	0.247
ARFIMA	0.207	0.16	0.242	0.107
Chronos-Bolt-S	0.168	0.12	0.174	0.277
Chronos-Bolt-B	**0.167**	**0.12**	0.176	**0.279**
Moirai-2.0-S	0.222	0.19	0.329	-0.008
Lag-Llama	0.207	0.17	0.305	0.080
Kronos	0.170	0.14	0.299	0.245
Panel: $h = 5$				
HAR	0.082	0.12	0.143	0.284
HAR-J	0.079	0.12	0.149	0.305
HAR-RS	0.098	0.13	0.240	0.173
HARQ	0.085	0.12	0.399	0.259
Log-HAR	0.080	0.11	**0.120**	0.306
ARFIMA	0.098	0.13	0.178	0.147
Chronos-Bolt-S	**0.075**	**0.10**	0.128	**0.353**
Chronos-Bolt-B	0.075	0.10	0.129	0.353
Moirai-2.0-S	0.110	0.17	0.250	-0.025
Lag-Llama	0.098	0.14	0.194	0.117
Kronos	0.076	0.12	0.198	0.311
Panel: $h = 22$				
HAR	0.052	0.13	0.186	0.171
HAR-J	0.052	0.12	0.197	0.177
HAR-RS	0.065	0.13	0.223	0.004
HARQ	0.053	0.12	0.367	0.155
Log-HAR	0.049	0.12	**0.147**	0.205
ARFIMA	0.060	0.12	0.185	0.033
Chronos-Bolt-S	**0.047**	**0.10**	0.161	**0.250**
Chronos-Bolt-B	0.047	0.10	0.163	0.248
Moirai-2.0-S	0.064	0.16	0.212	-0.093
Lag-Llama	0.053	0.13	0.164	0.119
Kronos	0.050	0.12	0.358	0.152

Forecast accuracy for 5 FX pairs (VOLARE). Values are averages across the five currency pairs. See Table 5.1 for variable definitions.

Model	MSE ($\times 10^{-6}$)	MAE ($\times 10^{-4}$)	QLIKE	R^2_{OOS}
Panel: $h = 1$				
HAR	14.818	2.72	0.225	0.244
HAR-J	14.958	2.65	0.707	**0.258**
HAR-RS	15.503	2.81	$5.82^{\dagger}$	0.180
HARQ	15.156	2.65	$22.34^{\dagger}$	0.089
Log-HAR	**14.143**	2.09	**0.193**	0.225
ARFIMA	16.437	2.95	0.346	0.009
Chronos-Bolt-S	14.335	1.99	0.213	0.237
Chronos-Bolt-B	14.251	**1.95**	0.221	0.246
Moirai-2.0-S	14.509	1.99	0.289	0.213
Lag-Llama	14.860	2.77	$1.12^{\dagger}$	0.074
Kronos	14.387	2.31	$10.57^{\dagger}$	-2.099
Panel: $h = 5$				
HAR	4.574	2.56	0.434	0.253
HAR-J	4.918	2.48	0.839	0.245
HAR-RS	4.615	2.60	$4.07^{\dagger}$	0.161
HARQ	4.792	2.49	$22.10^{\dagger}$	0.181
Log-HAR	4.442	1.98	**0.165**	0.226
ARFIMA	6.262	2.70	0.269	-0.040
Chronos-Bolt-S	**4.352**	1.80	0.185	0.273
Chronos-Bolt-B	4.384	**1.77**	0.196	**0.275**
Moirai-2.0-S	4.523	1.80	0.207	0.246
Lag-Llama	4.859	2.48	0.333	0.072
Kronos	4.498	2.10	$10.23^{\dagger}$	-5.953
Panel: $h = 22$				
HAR	2.169	2.77	$1.08^{\dagger}$	0.078
HAR-J	2.234	2.73	$1.23^{\dagger}$	0.071
HAR-RS	2.201	2.89	$4.01^{\dagger}$	-0.034
HARQ	2.126	2.66	$17.88^{\dagger}$	-0.051
Log-HAR	2.001	2.16	0.280	0.130
ARFIMA	3.681	2.65	0.286	-0.296
Chronos-Bolt-S	1.873	1.92	0.282	0.168
Chronos-Bolt-B	1.886	1.87	0.286	0.167
Moirai-2.0-S	**1.855**	**1.83**	**0.238**	**0.201**
Lag-Llama	2.062	2.29	0.336	0.082
Kronos	2.031	2.22	$12.89^{\dagger}$	-0.726

Forecast accuracy for 5 futures contracts (VOLARE). Values are averages across the five contracts. See Table 5.1 for variable definitions.

Foreign exchange. The FX sample covers five major currency pairs with lower volatility levels ($\bar{RV} \approx 3 \times 10^{-5}$) and moderate persistence ($\bar{\rho}_1 = 0.52$). For FX, Chronos-Bolt

consistently outperforms Moirai 2.0. At $h = 1$, Log-HAR achieves the lowest QLIKE (0.163), with Chronos-Bolt-Small close behind (0.174), while Moirai 2.0 underperforms substantially (QLIKE = 0.329, $R^2_{\text{OOS}} = -0.008$). At $h = 5$, Log-HAR retains the QLIKE lead (0.120 vs. 0.128 for Chronos-Bolt-Small), but Chronos-Bolt achieves the best MSE and R^2_{OOS} (0.353). At $h = 22$, Log-HAR again leads on QLIKE (0.147), while Chronos-Bolt achieves the best R^2_{OOS} (0.250). Moirai 2.0's poor FX performance ($R^2_{\text{OOS}} < 0$ at all horizons) stands in contrast to its equity dominance and suggests that Moirai's training data or architecture is less suited to the smoother, lower-amplitude dynamics of currency volatility.

Futures. The futures sample displays the widest cross-asset heterogeneity in volatility dynamics. The E-mini S&P 500 (ES) has first-order autocorrelation of 0.779, while Gold (GC) shows near-zero autocorrelation ($\rho_1 = -0.001$). At $h = 1$, Log-HAR achieves the lowest QLIKE (0.193), followed by Chronos-Bolt-Small (0.213) and HAR (0.225). Moirai 2.0 ranks behind (0.289). At $h = 5$, Log-HAR retains the QLIKE lead (0.165), with Chronos-Bolt-Small at 0.185 and Chronos-Bolt-Base at 0.196. The TSFMs achieve the best MSE and R^2_{OOS}. At the monthly horizon, Moirai 2.0 recovers and achieves the lowest QLIKE (0.238), MSE, and MAE, while levels-based HAR variants produce inflated QLIKE (HAR: 1.08, HARQ: 17.88). Kronos exhibits strongly negative R^2_{OOS} for futures (as low as -0.60 for Natural Gas), consistent with its small parameter count (8M) and its OHLCV-centric pretraining being poorly matched to the synthetic input we construct from univariate RV.

Cross-asset patterns. Three patterns emerge from the cross-asset comparison. First, Log-HAR is consistently the best or near-best econometric specification across all asset classes, confirming the practical value of the log transformation for ensuring forecast positivity and aligning with QLIKE's loss structure. Second, the relative ranking of TSFMs shifts across asset classes: Moirai 2.0 dominates equities (where persistence is high), while Chronos-Bolt performs better for FX and short-horizon futures (where dynamics are smoother or less persistent). Third, all levels-based HAR extensions produce inflated QLIKE for futures,

where extreme volatility events in Natural Gas and Crude Oil expose the negative-forecast problem. These asset-class-dependent results highlight the importance of multi-market evaluations: a model that appears dominant on equities alone may not generalize to other asset classes.

5.3 Statistical Significance and Model Confidence Sets

Having established the aggregate performance rankings, we now test the statistical significance of these differences. We apply four complementary tools: Model Confidence Sets to identify the subset of models that cannot be distinguished from the best, pairwise Diebold–Mariano tests (Diebold and Mariano, 1995) to quantify directional win rates, Mincer–Zarnowitz regressions (Mincer and Zarnowitz, 1969) to assess forecast efficiency, and Giacomini–Rossi fluctuation tests (Giacomini and Rossi, 2010) to detect time variation in relative performance.

Table 5.3 reports Model Confidence Set inclusion rates for 40 equities, and Table 5.3 summarizes pairwise Diebold–Mariano test results across all 50 assets.

Model	$h = 1$ (%)	$h = 5$ (%)	$h = 22$ (%)
HAR	42.5	42.5	20.0
HAR-J	17.5	35.0	20.0
HAR-RS	2.5	2.5	12.5
HARQ	0.0	0.0	7.5
Log-HAR	90.0	65.0	75.0
ARFIMA	2.5	30.0	87.5
Chronos-Bolt-S	42.5	30.0	37.5
Chronos-Bolt-B	37.5	25.0	30.0
Moirai-2.0-S	77.5	100.0	97.5
Lag-Llama	0.0	5.0	95.0
Kronos	2.5	0.0	7.5

Model Confidence Set inclusion rates for 40 U.S. equities (VOLARE). Each cell reports the percentage of stocks for which the model is included in the MCS at the 10% significance level (QLIKE loss, $T_{\max}$ statistic, block bootstrap with $B = 10{,}000$).

Model	$h = 1$	$h = 5$	$h = 22$
HAR	40.8%	23.0%	7.4%
HAR-J	24.0%	20.2%	7.0%
HAR-RS	13.8%	4.0%	2.0%
HARQ	0.6%	1.0%	0.4%
Log-HAR	66.8%	41.6%	22.4%
ARFIMA	30.0%	17.8%	19.4%
Chronos-Bolt-S	51.8%	36.4%	17.8%
Chronos-Bolt-B	51.2%	32.8%	15.8%
Moirai-2.0-S	54.4%	52.4%	29.2%
Lag-Llama	16.8%	13.2%	22.0%
Kronos	13.0%	8.6%	1.6%

Diebold–Mariano test: pairwise win rates. Each cell reports the percentage of pairwise DM tests (across 50 assets $\times$ 10 opponents = 500 tests) in which the row model has significantly lower QLIKE at the 5% level.

MCS results for equities. At $h = 1$, Log-HAR enters the MCS for 90% of assets, confirming its status as the single most reliable specification at the daily horizon. Moirai 2.0 enters for 77.5%, Chronos-Bolt-Small and HAR each for 42.5%. Among the additional TSFMs, Kronos enters for only 2.5% and Lag-Llama for 0%, indicating that these models are systematically outperformed at short horizons. Kronos's weak performance likely reflects its small parameter count (approximately 8M), which limits capacity relative to Chronos-Bolt (200M) and Moirai 2.0 (305M). Levels-based HAR variants are also largely excluded: HAR-RS enters for 2.5% and HARQ for 0%.

At $h = 5$, the picture shifts. Moirai 2.0 enters the MCS for 100% of assets, followed by Log-HAR (65%), HAR (42.5%), and Chronos-Bolt-Small (30%). Lag-Llama enters for only 5% and Kronos for 0%.

At $h = 22$, the ranking changes markedly. Moirai 2.0 enters for 97.5%, Lag-Llama for 95%, ARFIMA for 87.5%, Log-HAR for 75%, and Chronos-Bolt-Small for 37.5%. Lag-Llama's sharp improvement from 0% at $h = 1$ to 95% at $h = 22$ suggests that its lag-based tokenization captures long-range dependencies more effectively than short-term dynamics. Kronos remains weak (7.5%). The high ARFIMA inclusion rate at $h = 22$ is consistent with

its long-memory parameterization: at the monthly horizon, fractional differencing captures slow mean reversion that the HAR's three fixed lags approximate less precisely.

Diebold–Mariano pairwise tests. The DM test is computed pairwise: for each of the 11 models, we test it against each of the remaining 10 models on each of the 50 assets, yielding $11 \times 10 \times 50 = 5{,}500$ individual tests. Table 5.3 reports the fraction of pairwise DM tests in which each model achieves significantly lower QLIKE at the 5% level, aggregated across 50 assets and 10 opponents. At $h = 1$, Log-HAR wins the highest fraction of pairwise comparisons (66.8%), followed by Moirai 2.0 (54.4%), Chronos-Bolt-Small (51.8%), and HAR (40.8%). The TSFMs' strong showing on this directional metric (where each win requires both statistical significance and lower loss) confirms that their QLIKE advantage over levels-based HAR variants is not driven by a few outlier assets. At $h = 5$ and $h = 22$, Moirai 2.0 achieves the highest win rate (52.4% and 29.2%, respectively), reflecting its broad dominance at longer horizons. Lag-Llama's win rate rises from 16.8% at $h = 1$ to 22.0% at $h = 22$, consistent with its improved MCS inclusion at longer horizons.

Forecast efficiency: Mincer–Zarnowitz regressions. Table 5.3 reports average Mincer–Zarnowitz regression results across 50 assets. The MZ regression, $RV_t = \alpha + \beta \hat{f}_t + \varepsilon_t$, tests whether forecasts are efficient: under the null, $\alpha = 0$ and $\beta = 1$.

At $h = 1$, Moirai 2.0 and ARFIMA produce the most well-calibrated forecasts ($\hat{\beta} = 1.05$ for both), closest to the theoretical ideal of unity. The HAR family overshoots modestly ($\hat{\beta} = 1.29$ for HAR, 1.25 for HAR-J), while Log-HAR overshoots substantially ($\hat{\beta} = 1.77$), reflecting the bias-correction step that converts log-space forecasts back to levels. Among TSFMs, Chronos-Bolt achieves $\hat{\beta} = 1.20$ and the highest MZ R^2 (0.40–0.42), indicating that its forecasts explain a larger fraction of the variation in realized values than any econometric benchmark. The joint efficiency null ($\alpha = 0$, $\beta = 1$) is rejected for 98–100% of assets for all TSFMs, compared to 54–58% for HAR and HAR-J. This seemingly paradoxical result (TSFMs are more accurate on MSE and R^2 but more frequently reject efficiency) arises

because the MZ test has greater power when R^2 is high. A systematic $\hat{\beta} > 1$ indicates that TSFMs' forecast variation is slightly compressed relative to realized variation, a pattern consistent with the shrinkage implicit in their pretrained priors.

At $h = 22$ (Table A), all econometric models exhibit substantial undershoot ($\hat{\beta} = 0.30$–0.58), with MZ R^2 falling to 0.07–0.14. TSFMs retain substantially better calibration: Moirai 2.0 achieves $\hat{\beta} = 0.99$ and the highest R^2 (0.19), while Chronos-Bolt-Small achieves $\hat{\beta} = 1.12$. This confirms that TSFMs' advantage at longer horizons is not only in loss-function terms but also in the structural quality of their forecasts.

Model	$\hat{\alpha}$	$\hat{\beta}$	R^2	F-stat	p(joint)	% Reject
Econometric benchmarks						
HAR	−0.0000	1.288	0.365	20.69	0.154	58%
HAR-J	−0.0000	1.245	0.374	11.02	0.146	54%
HAR-RS	0.0000	1.123	0.309	24.11	0.099	72%
HARQ	0.0001	0.893	0.142	13.52	0.028	82%
Log-HAR	−0.0001	1.770	0.308	8.25	0.022	90%
ARFIMA	0.0000	1.054	0.171	18.93	0.059	84%
Foundation models						
Chronos-Bolt-S	0.0000	1.200	0.402	34.50	0.004	98%
Chronos-Bolt-B	0.0000	1.253	0.423	42.74	0.003	98%
Moirai-2.0-S	0.0000	1.050	0.378	22.17	0.002	98%
Lag-Llama	−0.0001	1.499	0.147	7.21	0.041	76%
Kronos	0.0000	1.204	0.354	103.03	0.000	100%

Mincer–Zarnowitz forecast efficiency regression, $h = 1$. Cross-sectional averages across 50 assets (40 equities, 5 FX, 5 futures). The regression is $RV_t = \alpha + \beta \hat{f}_t + \varepsilon_t$; under forecast efficiency, $\alpha = 0$ and $\beta = 1$. The F-statistic tests the joint null with HAC (Newey–West) standard errors. % Reject reports the fraction of assets for which the joint null is rejected at the 5% level.

Time-varying performance: Giacomini–Rossi fluctuation test. Figure 5 plots the rolling DM statistic (QLIKE loss) of each model against HAR at $h = 1$, averaged across 50 assets. A negative rolling DM indicates that the comparison model produces lower QLIKE than HAR during that window; values beyond the ± 2.80 critical-value bands indicate rejection of the equal predictive ability null at the 5% level.

Three patterns emerge. First, the relative performance of all models against HAR is time-varying: the null of constant relative predictive ability is rejected for 100% of assets

for Log-HAR, ARFIMA, and Lag-Llama, and for 62–98% for the other TSFMs. This means that unconditional DM tests and MCS results, while informative about average performance, mask substantial temporal variation. Second, Log-HAR's advantage over HAR is concentrated in the 2019–2021 period, where it crosses above the upper critical value, indicating that it significantly *underperforms* HAR during the COVID-induced volatility spike. This is consistent with the log-space bias correction amplifying forecast errors during regime changes. After 2022, Log-HAR's advantage narrows and approaches zero. Third, TSFMs (which enter the sample around 2022 when sufficient context becomes available) immediately settle below the lower critical-value band, indicating persistent and significant outperformance of HAR. Moirai 2.0 shows the most stable advantage, while Kronos exhibits the most volatile path.

At $h = 22$, the fluctuation test rejects less frequently (24–90% across models), consistent with the well-known finding that forecast differences diminish at longer horizons. Kronos is the only TSFM that maintains a high rejection rate (90%) at the monthly horizon, driven by systematically different forecasts rather than superior accuracy.

Interpretation. The MCS, DM, MZ, and fluctuation test results jointly provide a multidimensional picture of relative forecast performance. The MCS and DM tests establish that the best zero-shot TSFMs have achieved statistical parity with the best econometric benchmarks. The MZ regressions reveal that TSFMs produce better-calibrated forecasts at long horizons ($\hat{\beta}$ closer to 1, higher R^2), while econometric models suffer severe undershoot at $h = 22$. The Giacomini–Rossi test shows that relative performance is not constant: Log-HAR's QLIKE advantage at $h = 1$ erodes over time, and TSFMs consistently outperform HAR from their entry into the sample onward.

At $h = 1$, Log-HAR remains the marginal leader on QLIKE, but Moirai 2.0 and Chronos-Bolt are statistically indistinguishable from it for most assets. At $h = 5$ and $h = 22$, Moirai 2.0 is the single model most consistently in the confidence set. However, the results also reveal substantial heterogeneity across TSFMs: Lag-Llama and Kronos are systemat-

ically excluded at short horizons, with Lag-Llama recovering only at $h = 22$ and Kronos performing poorly throughout. This heterogeneity suggests that architecture and pretraining data matter as much as the zero-shot paradigm itself. The practical implication is that TSFMs offer a viable alternative to HAR-family models, particularly for horizons beyond one day, but practitioners should evaluate multiple TSFM architectures rather than assuming uniform performance.

5.4 Portfolio Application: Realized Covariance Forecasting

The preceding analysis focused on univariate realized variance forecasting. We now extend to the multivariate setting, examining whether TSFM forecasts of realized covariance matrices can improve portfolio allocation. VOLARE provides daily realized covariance matrices for each asset class (stocks, FX, futures), computed from 5-minute intraday returns using the same methodology as the univariate realized variance measures described in Section 3. These covariance matrices serve as both the forecasting target and the ex-post evaluation benchmark.

The global minimum variance (GMV) portfolio solves the optimization problem

$$\mathbf{w}_t^{\text{GMV}} = \arg\min_{\mathbf{w}} \ \mathbf{w}'\hat{\Sigma}_{t+h|t}\,\mathbf{w} \quad \text{s.t.} \quad \mathbf{w}'\mathbf{1} = 1, \tag{9}$$

where $\hat{\Sigma}_{t+h|t}$ is the forecast covariance matrix at horizon h formed at time t, and $\mathbf{w}$ is the vector of portfolio weights constrained to sum to one (short-selling is permitted). The closed-form solution is $\mathbf{w}_t^{\text{GMV}} = \hat{\Sigma}_{t+h|t}^{-1}\mathbf{1}/(\mathbf{1}'\hat{\Sigma}_{t+h|t}^{-1}\mathbf{1})$. The GMV portfolio requires no expected return inputs and depends entirely on the covariance forecast, making it the natural vehicle for evaluating forecast quality: differences in realized portfolio variance map directly to differences in covariance forecast accuracy.

Covariance forecasting methods. We consider four covariance forecasting approaches. (i) *Element-HAR* (Chiriac and Voev, 2011): a separate HAR model is fitted to each element of

the realized covariance matrix, treating variances and covariances as independent univariate series. (ii) *HAR-DRD*: the realized covariance matrix is decomposed as $\Sigma_t = D_t R_t D_t$, where D_t is the diagonal matrix of realized volatilities and R_t is the realized correlation matrix; HAR models are fitted separately to each element of D_t and R_t, and the covariance forecast is reconstructed. This DRD decomposition ensures positive semi-definiteness by construction when the forecast correlation matrix is adjusted to be positive semi-definite (PSD). (iii–iv) *Chronos-Bolt-Small* and *Moirai 2.0*: zero-shot element-wise TSFM forecasts applied to each covariance pair, followed by PSD projection. We also include the $1/N$ equal-weight portfolio as a naive benchmark.

Table 5.4 reports the average realized variance of GMV portfolios constructed from each forecasting model, along with portfolio turnover and concentration metrics. We present results in the same order as the univariate analysis: equities, FX, and futures.

For the 40-stock equity portfolio, the dimensionality problem becomes acute. Only HAR-DRD produces sensible portfolios, achieving a realized variance of 5.11×10^{-5} at $h = 1$ (a 29.4% reduction relative to $1/N$, 7.24×10^{-5}). All element-wise approaches (Element-HAR, Chronos-Bolt, Moirai 2.0) fail catastrophically: their realized variances exceed $1/N$ by two to three orders of magnitude, with average daily turnover above 7 (Chronos-Bolt) and 50 (Moirai). The problem is that element-wise forecasting of 820 unique covariance pairs produces assembled matrices with extremely small eigenvalues. PSD projection with a floor of 10^{-10} yields near-singular matrices whose inverse concentrates portfolio weight on a few assets. The DRD decomposition avoids this by separately modeling volatilities and correlations, ensuring that the reconstructed matrix inherits the PSD structure of the correlation forecast. This result highlights a practical limitation: element-wise TSFM covariance forecasting does not scale to high-dimensional portfolios without structural constraints. To verify that this failure is not an artifact of the eigenvalue floor, we re-projected all stocks forecast

matrices using a floor of 10^{-6} (six orders of magnitude larger than the default 10^{-10}) and re-evaluated the GMV portfolios. The higher floor reduces turnover substantially (from 13.7 to 2.2 for Chronos-Bolt and from 52.7 to 7.2 for Moirai at $h = 1$) but realized variances remain 5–7 times above $1/N$ (3.6×10^{-4} and 5.0×10^{-4} vs. 7.2×10^{-5}). HAR-DRD continues to outperform $1/N$ (4.6×10^{-5}). The failure of element-wise approaches thus reflects a fundamental structural problem, not a tuning parameter.

For the 5-asset FX portfolio, all model-based approaches achieve substantial variance reduction relative to $1/N$. At $h = 1$, Chronos-Bolt-Small achieves the lowest average realized variance (3.01×10^{-6}), a 44.9% reduction relative to $1/N$ (5.46×10^{-6}). HAR-DRD is close behind (3.11×10^{-6}, 43.0% reduction), while Element-HAR (3.52×10^{-6}, 35.5%) and Moirai 2.0 (3.55×10^{-6}, 35.0%) are somewhat less effective. At $h = 5$, Chronos-Bolt and HAR-DRD maintain their advantage (45.5% and 44.7% reduction, respectively), while Moirai 2.0 produces an unstable covariance forecast that results in higher realized variance than $1/N$. At $h = 22$, all methods except Moirai 2.0 produce comparable variance reductions (31–39%).

For the 5-contract futures portfolio, Chronos-Bolt-Small achieves the lowest realized variance at $h = 1$ (3.36×10^{-5}, a 68.8% reduction vs. $1/N$) and $h = 5$ (3.55×10^{-5}, 67.2%). Moirai 2.0 performs worse, particularly at $h = 5$, where its realized variance exceeds $1/N$ by a factor of 3.5. This failure stems from Moirai's element-wise covariance forecasts being internally inconsistent: the assembled matrices are not positive semi-definite, and the eigenvalue floor used in PSD projection (10^{-10}) is too small relative to the matrix scale (typical eigenvalues $\sim 10^{-5}$). The resulting near-singular matrices produce extreme portfolio weights on a handful of dates, which dominate the average realized variance. Chronos-Bolt does not suffer this problem because its forecasts produce matrices with minimum eigenvalues around

10^{-6}, keeping condition numbers below 10.

Turnover. The $1/N$ portfolio has zero turnover by construction. Among model-based portfolios in the low-dimensional FX and futures settings, HAR-DRD achieves the lowest turnover (0.080 for FX at $h = 1$), followed by Moirai 2.0 (0.060) and Chronos-Bolt (0.116). The DRD decomposition's smoothness advantage comes from its separate modeling of volatilities and correlations, which dampens spurious variation in off-diagonal elements.

Covariance forecast accuracy. Table 5.4 reports element-wise forecast accuracy across all three asset classes. Two metrics are computed: the normalized Frobenius norm of the forecast error matrix ($\|\hat{\Sigma}_t - \Sigma_t\|_F/N$) and the QLIKE loss on diagonal (variance) elements. The results show that TSFMs produce competitive covariance forecasts at the element level, even when the assembled matrices are ill-conditioned for portfolio optimization. For FX at $h = 1$, Chronos-Bolt achieves the lowest Frobenius norm (0.86×10^{-5}) and near-best diagonal QLIKE (0.145 vs. 0.133 for HAR-DRD). For futures, Chronos-Bolt and HAR-DRD are comparable on Frobenius norm, with HAR-DRD leading on QLIKE at $h = 5$ and $h = 22$. For stocks, Chronos-Bolt achieves the lowest Frobenius norm at $h = 1$ and $h = 5$, while HAR-DRD leads on diagonal QLIKE. Moirai 2.0 is competitive on stocks (QLIKE 0.175 at $h = 1$, matching Chronos-Bolt) but weaker on FX (QLIKE 0.308). The disconnect between element-wise accuracy and portfolio performance (Chronos-Bolt's Frobenius norm is lower than HAR-DRD's for stocks, yet its portfolios fail) confirms that the problem lies in the spectral properties of the assembled matrix, not in the accuracy of individual element forecasts. Potential solutions include applying the DRD decomposition to TSFM forecasts (forecasting volatilities and correlations separately, then reconstructing), imposing shrinkage toward a structured target (Ledoit and Wolf, 2004), or using natively multivariate TSFM architectures that produce positive semi-definite outputs by construction.

Model	h	Avg RV	Std RV	Turnover	Max Wt
Stocks (40 assets)					
$1/N$	1	72.38	200.30	—	—
HAR-DRD		**51.10**	124.55	0.788	0.159
Chronos-Bolt-S		$> 10^4$	—	13.695	—
Element-HAR		$> 10^4$	—	12.672	—
Moirai-2.0-S		$> 10^4$	—	52.690	—
$1/N$	5	72.39	130.35	—	—
HAR-DRD		**52.81**	81.06	0.789	0.159
Chronos-Bolt-S		$> 10^4$	—	7.639	—
Element-HAR		$> 10^4$	—	12.691	—
Moirai-2.0-S		$> 10^4$	—	48.779	—
$1/N$	22	72.21	93.25	—	—
HAR-DRD		**57.94**	80.69	0.790	0.159
Chronos-Bolt-S		$> 10^4$	—	7.710	—
Element-HAR		$> 10^4$	—	12.774	—
Moirai-2.0-S		$> 10^4$	—	51.283	—
Forex (5 assets)					
$1/N$	1	5.46	10.30	—	—
Element-HAR		3.52	6.50	0.096	0.383
HAR-DRD		3.11	4.50	0.080	0.370
Chronos-Bolt-S		**3.01**	4.78	0.116	0.369
Moirai-2.0-S		3.55	9.83	0.060	0.384
$1/N$	5	5.73	6.48	—	—
Element-HAR		3.58	5.45	0.096	0.383
HAR-DRD		3.17	2.99	0.080	0.370
Chronos-Bolt-S		**3.12**	3.69	0.090	0.369
Moirai-2.0-S		21.72	1083.82	0.104	0.397
$1/N$	22	5.43	4.32	—	—
Element-HAR		3.71	5.89	0.096	0.383
HAR-DRD		3.30	2.45	0.080	0.370
Chronos-Bolt-S		**3.29**	3.00	0.074	0.370
Moirai-2.0-S		3.52	6.06	0.046	0.382
Futures (5 contracts)					
$1/N$	1	107.77	316.84	—	—
Element-HAR		37.47	47.56	0.116	0.518
HAR-DRD		35.31	35.42	0.129	0.523
Chronos-Bolt-S		**33.63**	33.30	0.114	0.520
Moirai-2.0-S		45.08	156.03	0.264	0.552
$1/N$	5	107.88	207.81	—	—
Element-HAR		39.89	83.62	0.116	0.518
HAR-DRD		36.83	33.59	0.129	0.523
Chronos-Bolt-S		**35.46**	32.45	0.088	0.514
Moirai-2.0-S		376.26	19571.31	0.321	0.560
$1/N$	22	108.35	151.33	—	—
Element-HAR		42.66	50.69	0.116	0.518
HAR-DRD		40.75	38.02	0.129	0.523
Chronos-Bolt-S		**39.88**	44.59	0.065	0.503
Moirai-2.0-S		59.63	670.30	0.154	0.515

Global minimum variance portfolio performance. Avg RV is the average daily realized variance of the GMV portfolio ($\times 10^{-6}$). Turnover is the average daily absolute weight change. For stocks, element-wise approaches (Element-HAR, Chronos-Bolt, Moirai) produce near-singular matrices that cause extreme portfolio weights; their Avg RV values are orders of magnitude above $1/N$ and are omitted from the Std RV and Max Wt columns. Bold indicates the lowest Avg RV in each panel.

	Model	$h=1$		$h=5$		$h=22$	
		Frob ($\times 10^{-5}$)	QLIKE	Frob ($\times 10^{-5}$)	QLIKE	Frob ($\times 10^{-5}$)	QLIKE
Forex	Element-HAR	1.06	0.149	0.89	0.105	0.88	0.116
	HAR-DRD	0.96	0.133	0.79	**0.092**	0.79	**0.112**
	Chronos-Bolt-S	**0.86**	**0.145**	**0.72**	0.112	**0.72**	0.151
	Moirai-2.0-S	1.38	0.308	1.28	0.245	1.17	0.211
Futures	Element-HAR	21.64	0.227	20.04	0.183	20.46	0.214
	HAR-DRD	15.06	0.158	**13.35**	**0.125**	**13.71**	**0.175**
	Chronos-Bolt-S	**14.87**	**0.173**	13.74	0.160	14.87	0.260
	Moirai-2.0-S	15.20	0.228	13.88	0.181	14.41	0.217
Stocks	Element-HAR	7.22	0.196	6.18	0.171	6.46	0.212
	HAR-DRD	6.87	0.166	5.88	**0.150**	**5.81**	**0.204**
	Chronos-Bolt-S	**6.72**	**0.176**	**5.84**	0.181	5.89	0.303
	Moirai-2.0-S	7.48	0.175	6.37	0.160	6.20	0.237

Covariance forecast accuracy by asset class. Avg Frobenius is the average normalized Frobenius norm of the forecast error matrix ($\|\hat{\Sigma}_t - \Sigma_t\|_F/N$). Diag QLIKE is the average QLIKE computed on diagonal (variance) elements. Bold indicates the best value in each panel×horizon cell.

6 Robustness Checks

We assess the stability of our main findings along several dimensions.

Sub-sample stability: pre- and post-COVID. Table 6 splits the equity sample at March 1, 2020, and reports forecast accuracy separately for the pre-COVID (2015–2020) and post-COVID (2020–2026) periods. The sub-sample analysis reveals that TSFM performance is regime-robust, while levels-based econometric models are not.

In the pre-COVID period, Log-HAR achieves the lowest QLIKE at all horizons (0.206 at $h = 1$, 0.182 at $h = 5$, 0.321 at $h = 22$), though Chronos-Bolt-Base matches it on MSE and R^2_{OOS} at $h = 1$. The pre-COVID period features lower average volatility and fewer extreme events, conditions that favor the parsimonious HAR specification.

In the post-COVID period, TSFMs dominate across all metrics and horizons. Moirai 2.0 achieves the lowest QLIKE at $h = 1$ (0.213 vs. 0.225 for Log-HAR), $h = 5$ (0.174 vs. 0.252), and $h = 22$ (0.168 vs. 0.335), while also achieving the highest R^2_{OOS} (0.430, 0.378, 0.148). The levels-based HAR variants suffer extreme QLIKE inflation in the high-volatility regime: HARQ's QLIKE rises from 2.35 pre-COVID to 23.63 post-COVID at $h = 1$, because the near-zero forecast problem is exacerbated when realized values are large. Log-HAR is more

stable across regimes than its levels-based counterparts but still falls behind the TSFMs on all metrics in the post-COVID sample. The TSFM heterogeneity pattern persists across sub-samples: Lag-Llama underperforms in both periods at $h = 1$ (QLIKE 0.433 pre-COVID, 0.385 post-COVID) but improves at longer horizons, while Kronos achieves the lowest MSE pre-COVID at $h = 1$ (0.082) but has inflated QLIKE in both periods due to systematic forecast bias. These results suggest that TSFMs' advantage is concentrated in volatile market conditions, where their flexible nonparametric forecasts adapt more effectively than the HAR's linear structure.

Estimation window: 252 vs. 512 days. A potential concern is that TSFMs receive 512 days of context while econometric models use 252-day estimation windows, creating an information asymmetry that may favor TSFMs. To address this, we re-estimate all six econometric benchmarks with a 512-day rolling window, matching the TSFM context length. Table 6 reports the results at $h = 1$ for the 40-equity sample. Five of six models improve under the longer window: HAR R^2_{OOS} rises from 0.340 to 0.435, and levels-HAR QLIKE inflation decreases substantially (e.g., HAR-RS QLIKE drops from 3.48 to 0.57). The exception is ARFIMA, whose long-memory estimation degrades with a 512-day window (R^2_{OOS} falls from 0.132 to -0.332). Cross-model R^2_{OOS} rankings are highly stable (Spearman $\rho = 0.94$, $p < 0.01$). Importantly, even after granting econometric models the same 512-day information set, Chronos-Bolt and Moirai zero-shot QLIKE (0.19–0.23; Table 5.1) remain competitive with or below the best 512-day econometric QLIKE (Log-HAR at 0.20), confirming that the TSFM advantage is not an artifact of the longer context window.

QLIKE floor sensitivity. Our QLIKE implementation floors forecasts at 10^{-10} to prevent division by zero. Because levels-based HAR variants occasionally produce near-zero forecasts, this floor choice could affect model rankings. Table 6 reports cross-asset mean QLIKE at $h = 1$ under four floor values (10^{-4}, 10^{-6}, 10^{-8}, 10^{-10}). QLIKE values are identical for floors $\leq 10^{-6}$ (Spearman rank correlation $\rho = 1.0$ vs. 10^{-10}), confirming that the floor never

Model	Pre-COVID				Post-COVID			
	MSE ($\times 10^{-6}$)	MAE ($\times 10^{-4}$)	QLIKE	R^2_{OOS}	MSE ($\times 10^{-6}$)	MAE ($\times 10^{-4}$)	QLIKE	R^2_{OOS}
Panel A: $h = 1$								
HAR	0.089	0.984	0.494	0.223	0.328	1.684	1.809†	0.330
HAR-J	0.090	0.976	0.395	0.204	0.324	1.667	1.539†	0.345
HAR-RS	0.102	1.037	1.312†	0.158	0.356	1.797	5.151†	0.259
HARQ	0.091	0.959	2.346†	0.175	0.406	1.780	23.625†	0.091
Log-HAR	0.087	0.862	**0.206**	0.242	0.348	1.545	0.225	0.225
ARFIMA	0.123	1.148	0.328	-0.008	0.392	1.896	0.312	0.147
Chronos-Bolt-S	0.086	0.834	0.235	0.247	0.285	1.363	0.226	0.410
Chronos-Bolt-B	0.084	**0.823**	0.234	0.275	0.281	**1.353**	0.228	0.422
Moirai-2.0-S	0.084	0.857	0.251	0.235	**0.277**	1.360	**0.213**	**0.430**
Lag-Llama	0.106	1.188	0.433	0.007	0.388	1.899	0.385	0.120
Kronos	**0.082**	0.895	0.954	**0.287**	0.309	1.527	1.166†	0.340
Panel B: $h = 5$								
HAR	0.050	0.972	0.421	0.029	0.221	1.643	3.877†	0.193
HAR-J	0.052	0.966	0.427	0.018	0.214	1.621	3.529†	0.219
HAR-RS	0.061	1.025	2.051†	-0.098	0.249	1.777	9.622†	0.077
HARQ	0.050	0.954	1.814†	0.036	0.265	1.694	23.517†	-0.008
Log-HAR	**0.044**	0.839	**0.182**	0.190	0.226	1.500	0.252	0.108
ARFIMA	0.063	1.014	0.263	-0.063	0.227	1.672	0.237	0.167
Chronos-Bolt-S	0.047	0.791	0.241	0.168	0.170	1.256	0.200	0.348
Chronos-Bolt-B	0.046	**0.782**	0.240	0.187	0.172	1.261	0.204	0.337
Moirai-2.0-S	0.045	0.789	0.217	0.202	**0.162**	**1.216**	**0.174**	**0.378**
Lag-Llama	0.056	1.024	0.274	0.026	0.246	1.638	0.321	0.020
Kronos	0.045	0.842	0.700	**0.229**	0.195	1.421	0.953	0.220
Panel C: $h = 22$								
HAR	0.095	1.145	1.587†	-0.032	0.118	1.621	8.230†	-0.070
HAR-J	0.095	1.140	1.510†	-0.028	0.117	1.608	8.084†	-0.047
HAR-RS	0.101	1.197	3.796†	-0.084	0.133	1.739	11.923†	-0.243
HARQ	0.095	1.134	2.323†	-0.039	0.139	1.637	17.488†	-0.331
Log-HAR	**0.091**	1.062	**0.321**	-0.006	0.116	1.470	0.335	-0.102
ARFIMA	0.098	1.150	0.361	-0.065	0.119	1.418	0.205	-0.031
Chronos-Bolt-S	0.115	1.095	0.533	-0.029	0.092	1.152	0.211	0.121
Chronos-Bolt-B	0.115	1.090	0.531	-0.027	0.093	1.160	0.217	0.108
Moirai-2.0-S	0.113	**1.057**	0.427	**-0.002**	**0.089**	**1.084**	**0.168**	**0.148**
Lag-Llama	0.114	1.184	0.362	-0.007	0.105	1.285	0.239	-0.012
Kronos	0.117	1.178	1.016†	-0.029	0.102	1.298	0.854	0.003

Sub-sample forecast accuracy: pre-COVID (2015–2020) and post-COVID (2020–2026) periods for 40 U.S. equities (VOLARE). Bold indicates best in each column within each panel. † denotes QLIKE values inflated by near-zero forecasts.

	MSE ($\times 10^8$)		QLIKE		R^2_{OOS}	
Model	252	512	252	512	252	512
HAR	23.03	**21.38**	1.2715	**0.2107**	0.340	**0.435**
HAR-J	22.78	**20.89**	1.0356	**0.3590**	0.353	**0.442**
HAR-RS	25.18	**21.92**	3.4828	**0.5729**	0.272	**0.417**
HARQ	27.71	**24.64**	14.5773	**4.8653**	0.119	**0.270**
Log-HAR	23.98	**22.46**	0.2166	**0.1954**	0.249	**0.381**
ARFIMA	**28.39**	39.98	**0.3264**	0.4786	**0.132**	-0.332

Estimation window sensitivity: 252-day vs. 512-day rolling window ($h = 1$, 40 equities). Cross-asset mean metrics. Bold indicates the better window for each model.

binds at these levels. Only the aggressive 10^{-4} floor alters rankings ($\rho = 0.60$) by clipping levels-HAR forecasts that drive QLIKE inflation (e.g., HARQ drops from 16.04 to 0.38). Our default 10^{-10} floor is conservative and does not distort model comparisons.

Model	$1e-10$	$1e-08$	$1e-06$	$1e-04$
ARFIMA	0.3218	0.3218	0.3218	0.3195
Log-HAR	0.2189	0.2189	0.2189	0.2358
HAR-J	1.11	1.11	1.11	0.2461
HARQ	16.04	16.04	16.04	0.3827
HAR	1.38	1.38	1.38	0.2565
HAR-RS	3.63	3.63	3.63	0.2838
Lag-Llama	0.4018	0.4018	0.4018	0.3639
Chronos-Bolt-S	0.2292	0.2292	0.2292	0.2353
Kronos	1.09	1.09	1.09	0.2346
Moirai-2.0-S	0.2263	0.2263	0.2263	0.2274
Chronos-Bolt-B	0.2302	0.2302	0.2302	0.2351

QLIKE floor sensitivity: cross-asset mean QLIKE at $h = 1$ under four floor values. Rankings are stable across floors (Spearman $\rho > 0.99$).

A model's low loss could reflect good calibration (i.e., alignment between the scale of forecasts and realizations) rather than genuine predictive information about future volatility dynamics. The Mincer–Zarnowitz regression (Mincer and Zarnowitz, 1969) provides a natural correction: by regressing realized values on forecasts, one obtains intercept and slope

parameters that remove any linear bias. If a model's accuracy survives this correction, its forecasts contain information beyond what a simple affine rescaling could produce. To test this, we apply a recursive Mincer–Zarnowitz bias correction to all TSFM forecasts at $h = 1$: at each time t, we estimate $\hat{\alpha}_t$ and $\hat{\beta}_t$ from all observations up to $t-1$, then form the corrected forecast $\hat{f}_t^{\mathrm{MZ}} = \hat{\alpha}_t + \hat{\beta}_t f_t$ using an expanding estimation window (minimum 252 days). Table 6 compares original and MZ-corrected TSFM forecasts against econometric benchmarks, evaluated on the common post-warm-up sample. MZ correction improves QLIKE for the well-calibrated Chronos-Bolt models (Chronos-Bolt-Base: 0.2307 $\rightarrow$ 0.1918), but worsens QLIKE for models with systematic bias (Lag-Llama, Kronos) where the linear correction pushes some forecasts toward zero. MSE and R^2_{OOS} are largely unchanged, confirming that TSFM accuracy reflects genuine predictive content rather than calibration artifacts.

Model	Variant	MSE ($\times 10^8$)	QLIKE	R^2_{OOS}
HAR	Original	27.07	1.52	0.333
HAR-J	Original	26.76	1.22	0.347
HAR-RS	Original	29.44	4.04	0.266
HARQ	Original	32.98	17.97	0.105
Log-HAR	Original	28.41	0.2220	0.237
ARFIMA	Original	33.34	0.3279	0.120
Chronos-Bolt-S	Original	24.04	0.2298	0.403
Chronos-Bolt-S	MZ-Corrected	27.51	0.1969	0.322
Chronos-Bolt-B	Original	23.67	0.2307	0.415
Chronos-Bolt-B	MZ-Corrected	25.40	0.1918	0.383
Moirai-2.0-S	Original	23.37	0.2245	0.421
Moirai-2.0-S	MZ-Corrected	25.56	0.2238	0.331
Lag-Llama	Original	32.36	0.4162	0.117
Lag-Llama	MZ-Corrected	32.46	7.05	0.116
Kronos	Original	25.63	0.9452	0.342
Kronos	MZ-Corrected	25.88	7.84	0.348

MZ bias-corrected TSFM evaluation (cross-asset mean, $h = 1$). Each TSFM forecast $\hat{f}_t$ is replaced by $\hat{\alpha}+\hat{\beta}\hat{f}_t$ from a rolling Mincer–Zarnowitz regression, then re-evaluated on MSE, QLIKE, and R^2_{OOS}. Econometric benchmarks are shown unchanged for comparison.

7 Conclusion

This paper evaluates zero-shot time series foundation models for realized volatility and covariance forecasting across 50 assets spanning U.S. equities, foreign exchange, and futures. We compare four TSFMs (Chronos-Bolt, Moirai 2.0, Lag-Llama, and Kronos) against six econometric specifications at three forecast horizons using the VOLARE dataset, with formal statistical testing via Diebold–Mariano tests (Diebold and Mariano, 1995) and Model Confidence Sets (Hansen et al., 2011).

Three main findings emerge. First, the best zero-shot TSFMs are competitive with the best econometric benchmarks for univariate realized variance forecasting. This contrasts with prior negative findings on daily returns (Rahimikia et al., 2025) and on TimesFM 1.0 (Goel et al., 2025), and is consistent with the hypothesis that realized volatility's positive, persistent, mean-reverting dynamics align more closely with the series on which TSFMs were pretrained. However, not all TSFMs perform equally: Kronos is excluded from the MCS for virtually all assets, and Lag-Llama only enters at $h = 22$. Second, model rankings depend on loss function and horizon. Log-HAR achieves the lowest QLIKE at $h = 1$ for equities, while TSFMs dominate on MSE, MAE, and R^2_{OOS} at all horizons; at $h = 5$ and $h = 22$, Moirai 2.0 leads across all metrics. Third, rankings shift across asset classes: Chronos-Bolt outperforms on FX and futures, while Moirai 2.0 dominates equities.

We also evaluate realized covariance forecasting and its downstream application to global minimum variance portfolio construction. Element-wise zero-shot TSFM forecasts of realized covariance matrices achieve substantial variance reductions (35–69%) relative to equal-weight allocation in low-dimensional portfolios (5 assets), with Chronos-Bolt performing comparably to HAR-DRD baselines. However, element-wise TSFM covariance forecasting fails for the 40-stock portfolio, where the assembled matrices are ill-conditioned despite competitive element-level accuracy. Only the structurally constrained HAR-DRD decomposition produces viable high-dimensional allocations, highlighting that covariance forecasting requires structural constraints that current element-wise TSFM approaches lack.

Several limitations should be noted. We evaluate only zero-shot TSFM performance; parameter-efficient fine-tuning may yield further gains, though the adaptation mechanisms differ across architectures: Chronos-2 supports LoRA-based adaptation (Hu et al., 2022) through its encoder-decoder structure, while Lag-Llama uses epoch-based training through its GluonTS estimator interface, reflecting their distinct architectural designs. Our evaluation is restricted to point forecasts; density forecast evaluation (Adler et al., 2025) remains an open question. VOLARE provides alternative realized measures (subsampled realized variance, realized kernels, range-based estimators) whose effect on model rankings we leave for future work. Natively multivariate TSFM architectures (Ansari et al., 2025) or structural decompositions applied to TSFM outputs may address the high-dimensional covariance forecasting limitation.

The broader implication is that zero-shot time series foundation models have reached a level of accuracy that makes them a practical tool for volatility and covariance forecasting. For practitioners, Moirai 2.0 and Chronos-Bolt combine competitive accuracy with zero estimation overhead and immediate deployment to new assets. At the same time, the poor performance of some TSFMs at short horizons and the failure of element-wise covariance forecasting at high dimensions caution against treating TSFMs as a universal solution. Model selection and structural constraints remain essential.

References

Coen Adler, Yuxin Chang, Felix Draxler, Samar Abdi, and Padhraic Smyth. Beyond accuracy: Are time series foundation models well-calibrated? *arXiv preprint arXiv:2510.16060*, 2025.

Taha Aksu, Gerald Woo, Juncheng Liu, Xu Liu, Chenghao Liu, Silvio Savarese, Caiming

Xiong, and Doyen Sahoo. GIFT-Eval: A benchmark for general time series forecasting model evaluation. *arXiv preprint arXiv:2410.10393*, 2024.

Torben G. Andersen and Tim Bollerslev. Answering the skeptics: Yes, standard volatility models do provide accurate forecasts. *International Economic Review*, 39(4):885–905, 1998. doi: 10.2307/2527343.

Torben G. Andersen, Tim Bollerslev, Francis X. Diebold, and Paul Labys. The distribution of realized exchange rate volatility. *Journal of the American Statistical Association*, 96 (453):42–55, 2001. doi: 10.1198/016214501750332965.

Torben G. Andersen, Tim Bollerslev, Francis X. Diebold, and Paul Labys. Modeling and forecasting realized volatility. *Econometrica*, 71(2):579–625, 2003. doi: 10.1111/1468-0262.00418.

Torben G. Andersen, Tim Bollerslev, and Francis X. Diebold. Roughing it up: Including jump components in the measurement, modeling, and forecasting of return volatility. *Review of Economics and Statistics*, 89(4):701–720, 2007. doi: 10.1162/rest.89.4.701.

Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, Jasper Zschiegner, Danielle C. Maddix, Hao Wang, Michael W. Mahoney, Kari Torkkola, Andrew Gordon Wilson, Michael Bohlke-Schneider, and Yuyang Wang. Chronos: Learning the language of time series. *Transactions on Machine Learning Research*, 2024. arXiv:2403.07815.

Abdul Fatir Ansari, Oleksandr Shchur, Jaris Küken, Andreas Auer, Boran Han, Pedro Mercado, Syama Sundar Rangapuram, Huibin Shen, Lorenzo Stella, Xiyuan Zhang, Mononito Goswami, Shubham Kapoor, Danielle C. Maddix, Pablo Guerron, Tony Hu, Junming Yin, Nick Erickson, Prateek Mutalik Desai, Hao Wang, Huzefa Rangwala, George Karypis,

Yuyang Wang, and Michael Bohlke-Schneider. Chronos-2: From univariate to universal forecasting. *arXiv preprint arXiv:2510.15821*, 2025.

Richard T. Baillie, Tim Bollerslev, and Hans Ole Mikkelsen. Fractionally integrated generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 74(1):3–30, 1996. doi: 10.1016/S0304-4076(95)01749-6.

Ole E. Barndorff-Nielsen and Neil Shephard. Econometric analysis of realized volatility and its use in estimating stochastic volatility models. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 64(2):253–280, 2002. doi: 10.1111/1467-9868.00336.

Ole E. Barndorff-Nielsen and Neil Shephard. Power and bipower variation with stochastic volatility and jumps. *Journal of Financial Econometrics*, 2(1):1–37, 2004. doi: 10.1093/jjfinec/nbh001.

Ole E. Barndorff-Nielsen, Peter Reinhard Hansen, Asger Lunde, and Neil Shephard. Designing realized kernels to measure the ex post variation of equity prices in the presence of noise. *Econometrica*, 76(6):1481–1536, 2008. doi: 10.3982/ECTA6495.

Ole E. Barndorff-Nielsen, Silja Kinnebrock, and Neil Shephard. Measuring downside risk: Realised semivariance. In Tim Bollerslev, Jeffrey Russell, and Mark Watson, editors, *Volatility and Time Series Econometrics: Essays in Honor of Robert F. Engle*, pages 117–136. Oxford University Press, 2010.

Tim Bollerslev, Andrew J. Patton, and Rogier Quaedvlieg. Exploiting the errors: A simple approach for improved volatility forecasting. *Journal of Econometrics*, 192(1):1–18, 2016. doi: 10.1016/j.jeconom.2015.10.007.

Rafael Branco, Alexandre Rubesam, and Mauricio Zevallos. Forecasting realized volatility: Does anything beat linear models? *Journal of Empirical Finance*, 78:101524, 2024. doi: 10.1016/j.jempfin.2024.101524.

Alessio Brini and Giacomo Toscano. SpotV2Net: Multivariate intraday spot volatility forecasting via vol-of-vol-informed graph attention networks. *International Journal of Forecasting*, 41(3):1093–1111, 2025. doi: 10.1016/j.ijforecast.2024.11.004.

Andrea Bucci. Realized volatility forecasting with neural networks. *Journal of Financial Econometrics*, 18(3):502–531, 2020. doi: 10.1093/jjfinec/nbaa008.

Andrea Carriero, Davide Pettenuzzo, and Shubhranshu Shekhar. Macroeconomic forecasting with large language models. *arXiv preprint arXiv:2407.00890*, 2024.

Roxana Chiriac and Valeri Voev. Modelling and forecasting multivariate realized volatility. *Journal of Applied Econometrics*, 26(6):922–947, 2011. doi: 10.1002/jae.1152.

Kim Christensen, Mathias Siggaard, and Bezirgen Veliyev. A machine learning approach to volatility forecasting. *Journal of Financial Econometrics*, 21(5):1680–1727, 2023. doi: 10.1093/jjfinec/nbac020.

Fabrizio Cipollini, Giulia Cruciani, Giampiero M. Gallo, Alessandra Insana, Edoardo Otranto, and Fabio Spagnolo. VOLARE: VOLatility Archive for Realized Estimates. *arXiv preprint arXiv:2602.19732*, 2026.

Adam Clements and Daniel P. A. Preve. A practical guide to harnessing the HAR volatility model. *Journal of Banking & Finance*, 133:106285, 2021. doi: 10.1016/j.jbankfin.2021.106285.

Fulvio Corsi. A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2):174–196, 2009. doi: 10.1093/jjfinec/nbp001.

Abhimanyu Das, Weihao Kong, Rajat Sen, and Yichen Zhou. A decoder-only foundation model for time-series forecasting. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *Proceedings of Machine Learning Research*, pages 10148–10167. PMLR, 2024. arXiv:2310.10688.

Abhimanyu Das, Matthew Faw, Rajat Sen, and Yichen Zhou. In-context fine-tuning for time-series foundation models. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2025. arXiv:2410.24087.

Francis X. Diebold and Roberto S. Mariano. Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3):253–263, 1995. doi: 10.1080/07350015.1995.10524599.

M. Germán-Morales, A. J. Rivera-Rivas, M. J. del Jesus Díaz, and C. J. Carmona. Transfer learning with foundational models for time series forecasting using low-rank adaptations. *arXiv preprint arXiv:2410.11539*, 2024.

Raffaella Giacomini and Barbara Rossi. Forecast comparisons in unstable environments. *Journal of Applied Econometrics*, 25(4):595–620, 2010. doi: 10.1002/jae.1177.

Anubha Goel, Puneet Pasricha, and Juho Kanniainen. Time-series foundation AI model for value-at-risk forecasting. *arXiv preprint arXiv:2410.11773*, 2024.

Anubha Goel, Puneet Pasricha, Martin Magris, and Juho Kanniainen. Foundation time-series AI model for realized volatility forecasting. *arXiv preprint arXiv:2505.11163*, 2025.

Mononito Goswami, Konrad Szafer, Arjun Choudhry, Yifu Cai, Shuo Li, and Artur Dubrawski. MOMENT: A family of open time-series foundation models. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *Proceedings of Machine Learning Research*. PMLR, 2024. arXiv:2402.03885.

Elias Søvik Gunnarsson, Håkon Ramon Isern, Aristidis Kaloudis, Morten Risstad, Benjamin Vigdel, and Sjur Westgaard. Prediction of realized volatility and implied volatility indices using AI and machine learning: A review. *International Review of Financial Analysis*, 93: 103221, 2024. doi: 10.1016/j.irfa.2024.103221.

Divij Gupta, Anubhav Bhatti, and Surajsinh Parmar. Beyond LoRA: Exploring efficient fine-

tuning techniques for time series foundational models. *arXiv preprint arXiv:2409.11302*, 2024.

Peter Reinhard Hansen, Asger Lunde, and James M. Nason. The model confidence set. *Econometrica*, 79(2):453–497, 2011. doi: 10.3982/ECTA5771.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2022. arXiv:2106.09685.

Marwan Izzeldin, M. Kabir Hassan, Vasileios Pappas, and Mike Tsionas. Forecasting realised volatility using ARFIMA and HAR models. *Quantitative Finance*, 19(10):1627–1638, 2019. doi: 10.1080/14697688.2019.1600713.

Stanisław Łaniewski and Robert Ślepaczuk. Evaluating the Chronos foundation model for daily stock index forecasting. In *Proceedings of the International Conference on Information Systems Development (ISD 2025)*, 2025. doi: 10.62036/ISD.2025.48.

Olivier Ledoit and Michael Wolf. A well-conditioned estimator for large-dimensional covariance matrices. *Journal of Multivariate Analysis*, 88(2):365–411, 2004. doi: 10.1016/S0047-259X(03)00096-4.

Radmir Mishelevich Leushuis and Nicolai Petkov. Advances in forecasting realized volatility: A review of methodologies. *Financial Innovation*, 12:5, 2026. doi: 10.1186/s40854-025-00809-5.

Zhe Li, Xiangfei Qiu, Peng Chen, Yihang Wang, Hanyin Cheng, Yang Shu, Jilin Hu, Chenjuan Guo, Aoying Zhou, Christian S. Jensen, and Bin Yang. TSFM-Bench: A comprehensive and unified benchmark of foundation models for time series forecasting. *arXiv preprint arXiv:2410.11802*, 2024. KDD 2025.

Yuxuan Liang, Haomin Wen, Yuqi Nie, Yushan Jiang, Ming Jin, Dongjin Song, Shirui Pan, and Qingsong Wen. Foundation models for time series analysis: A tutorial and survey. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, 2024. arXiv:2403.14735.

Chenghao Liu, Taha Aksu, Juncheng Liu, Xu Liu, Hanshu Yan, Quang Pham, Silvio Savarese, Doyen Sahoo, Caiming Xiong, and Junnan Li. Moirai 2.0: When less is more for time series forecasting. *arXiv preprint arXiv:2511.11698*, 2025a.

Yong Liu, Haoran Zhang, Chenyu Li, Xiangdong Huang, Jianmin Wang, and Mingsheng Long. Timer: Generative pre-trained transformers are large time series models. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *Proceedings of Machine Learning Research*. PMLR, 2024. arXiv:2402.02368.

Yong Liu, Guo Qin, Zhiyuan Shi, Zhi Chen, Caiyin Yang, Xiangdong Huang, Jianmin Wang, and Mingsheng Long. Sundial: A family of highly capable time series foundation models. *arXiv preprint arXiv:2502.00816*, 2025b. ICML 2025 Oral (Top 1%).

Chuong Luong and Nikolai Dokuchaev. Forecasting of realised volatility with the random forests algorithm. *Journal of Risk and Financial Management*, 11(4):61, 2018. doi: 10.3390/jrfm11040061.

Ben A. Marconi. Time series foundation models for multivariate financial time series forecasting. *arXiv preprint arXiv:2507.07296*, 2025.

John A. Miller, Mohammed Aldosari, Farah Saeed, Nasid Habib Barna, Subas Rana, I. Budak Arpinar, and Ninghao Liu. A survey of deep learning and foundation models for time series forecasting. *arXiv preprint arXiv:2401.13912*, 2024.

Jacob A. Mincer and Victor Zarnowitz. The evaluation of economic forecasts. In Jacob A. Mincer, editor, *Economic Forecasts and Expectations: Analysis of Forecasting Behavior and Performance*, pages 1–46. National Bureau of Economic Research, New York, 1969.

Fernando Moreno-Pino and Stefan Zohren. DeepVol: Volatility forecasting from high-frequency data with dilated causal convolutions. *Quantitative Finance*, 24(8):1105–1127, 2024. doi: 10.1080/14697688.2024.2387222.

Quoc Anh Nguyen, Ce Guo, and Wayne Luk. Repurposing language models for FX volatility forecasting: A data-efficient and context-aware approach. In *Proceedings of the 6th ACM International Conference on AI in Finance (ICAIF '25)*. ACM, 2025. doi: 10.1145/3768292.3770386.

Andrew J. Patton. Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics*, 160(1):246–256, 2011. doi: 10.1016/j.jeconom.2010.03.034.

Andrew J. Patton and Kevin Sheppard. Good volatility, bad volatility: Signed jumps and the persistence of volatility. *Review of Economics and Statistics*, 97(3):683–697, 2015. doi: 10.1162/REST_a_00503.

Eghbal Rahimikia, Hao Ni, and Weiguan Wang. Re(Visiting) time series foundation models in finance. *arXiv preprint arXiv:2511.18578*, 2025.

Kashif Rasul, Arjun Ashok, Andrew Robert Williams, Hena Ghonia, Rishika Bhagwatkar, Arian Khorasani, Mohammad Javad Darvishi Bayazi, George Adamopoulos, Roland Riachi, Nadhir Hassen, Marin Biloš, Sahil Garg, Anderson Schneider, Nicolas Chapados, Alexandre Drouin, Valentina Zantedeschi, Yuriy Nevmyvaka, and Irina Rish. Lag-llama: Towards foundation models for probabilistic time series forecasting. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024. arXiv:2310.08278.

Oleksandr Shchur, Abdul Fatir Ansari, Caner Turkmen, Lorenzo Stella, Nick Erickson, Pablo Guerron, Michael Bohlke-Schneider, and Yuyang Wang. fev-bench: A realistic benchmark for time series forecasting. *arXiv preprint arXiv:2509.26468*, 2025.

Yu Shi, Zongliang Fu, Shuo Chen, Bohan Zhao, Wei Xu, Changshui Zhang, and Jian

Li. Kronos: A foundation model for the language of financial markets. *arXiv preprint arXiv:2508.02739*, 2025.

Sébastien Valeyre and Sofiane Aboura. LLMs for time series: An application for single stocks and statistical arbitrage. *arXiv preprint arXiv:2412.09394*, 2024.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, pages 5998–6008, 2017. arXiv:1706.03762.

Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Kronos-T5: Unified time series forecasting via tokenized pretraining. Salesforce AI Research, 2024a. HuggingFace: Salesforce/kronos-t5-tiny.

Gerald Woo, Chenghao Liu, Akshat Kumar, Caiming Xiong, Silvio Savarese, and Doyen Sahoo. Unified training of universal time series forecasting transformers. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *Proceedings of Machine Learning Research*, pages 53140–53164. PMLR, 2024b. arXiv:2402.02592. Oral presentation.

Jiexia Ye, Yongzi Yu, Weiqi Zhang, Le Wang, Jia Li, and Fugee Tsung. Empowering time series analysis with foundation models: A comprehensive survey. *arXiv preprint arXiv:2405.02358*, 2024.

Chao Zhang, Yihuang Zhang, Mihai Cucuringu, and Zhongmin Qian. Volatility forecasting with machine learning and intraday commonality. *Journal of Financial Econometrics*, 22 (2):492–530, 2024. doi: 10.1093/jjfinec/nbad005.

Chao Zhang, Xingyue Pu, Mihai Cucuringu, and Xiaowen Dong. Graph-based methods for forecasting realized covariances. *Journal of Financial Econometrics*, 23(2):nbae026, 2025. doi: 10.1093/jjfinec/nbae026.

Zhuohang Zhu, Haodong Chen, Qiang Qu, and Vera Chung. FinCast: A foundation model for financial time-series forecasting. *arXiv preprint arXiv:2508.19609*, 2025. CIKM 2025.

A Supplementary Tables

This appendix collects additional results referenced in the main text. Table A reports cross-sectional medians rather than means for the 40-equity sample. Table A provides a per-asset QLIKE breakdown at $h = 1$ across all 50 assets. Tables A and A present Mincer–Zarnowitz regressions at $h = 5$ and $h = 22$.

Model	MSE ($\times 10^{-6}$)	MAE ($\times 10^{-4}$)	QLIKE	R^2_{OOS}
Panel: $h = 1$				
HAR	0.116	1.14	0.221	0.352
HAR-J	0.113	1.14	0.567	0.365
HAR-RS	0.124	1.20	$1.37^{\dagger}$	0.318
HARQ	0.173	1.24	$11.12^{\dagger}$	0.066
Log-HAR	0.139	1.12	**0.208**	0.237
ARFIMA	0.147	1.37	0.297	0.163
Chronos-Bolt-S	0.096	0.97	0.212	0.436
Chronos-Bolt-B	0.094	**0.95**	0.213	**0.459**
Moirai-2.0-S	**0.091**	0.98	0.208	0.451
Lag-Llama	0.156	1.38	0.384	0.134
Kronos	0.112	1.08	0.413	0.379
Panel: $h = 5$				
HAR	0.096	1.13	0.269	0.253
HAR-J	0.094	1.12	0.384	0.271
HAR-RS	0.109	1.19	$3.59^{\dagger}$	0.179
HARQ	0.126	1.17	$13.39^{\dagger}$	0.017
Log-HAR	0.111	1.08	0.205	0.133
ARFIMA	0.098	1.19	0.232	0.192
Chronos-Bolt-S	0.077	0.90	0.196	0.366
Chronos-Bolt-B	0.080	0.91	0.196	0.359
Moirai-2.0-S	**0.073**	**0.86**	**0.172**	**0.383**
Lag-Llama	0.128	1.20	0.295	0.039
Kronos	0.099	1.02	0.295	0.246
Panel: $h = 22$				
HAR	0.078	1.23	$3.40^{\dagger}$	0.007
HAR-J	0.076	1.22	$3.39^{\dagger}$	0.016
HAR-RS	0.082	1.30	$5.23^{\dagger}$	-0.090
HARQ	0.089	1.25	$9.29^{\dagger}$	-0.121
Log-HAR	0.081	1.13	0.296	-0.043
ARFIMA	0.075	1.10	0.269	0.073
Chronos-Bolt-S	0.071	0.94	0.319	0.085
Chronos-Bolt-B	0.072	0.97	0.326	0.082
Moirai-2.0-S	**0.069**	**0.89**	**0.248**	**0.112**
Lag-Llama	0.079	1.03	0.267	0.002
Kronos	0.078	1.06	0.542	0.024

Forecast accuracy for 40 U.S. equities (VOLARE), cross-sectional medians. Median aggregation is robust to outlier assets with degenerate forecasts. Bold indicates the best value in each column within each panel.

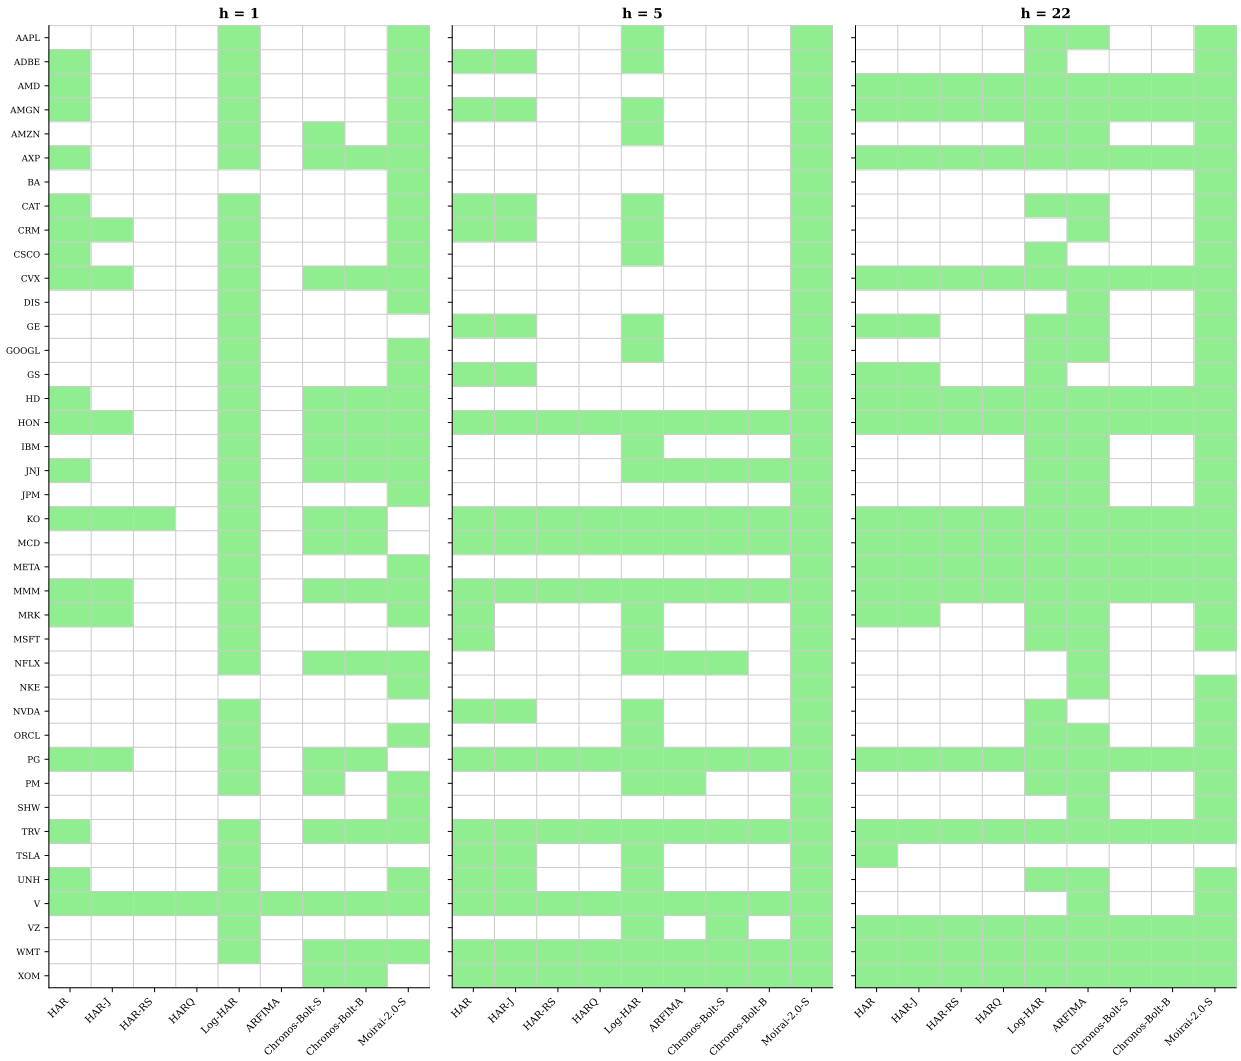

Figure 3: Model Confidence Set inclusion for 40 U.S. equities. Green cells indicate that the model is included in the MCS at the 10% significance level (QLIKE loss). Left: $h = 1$; center: $h = 5$; right: $h = 22$.

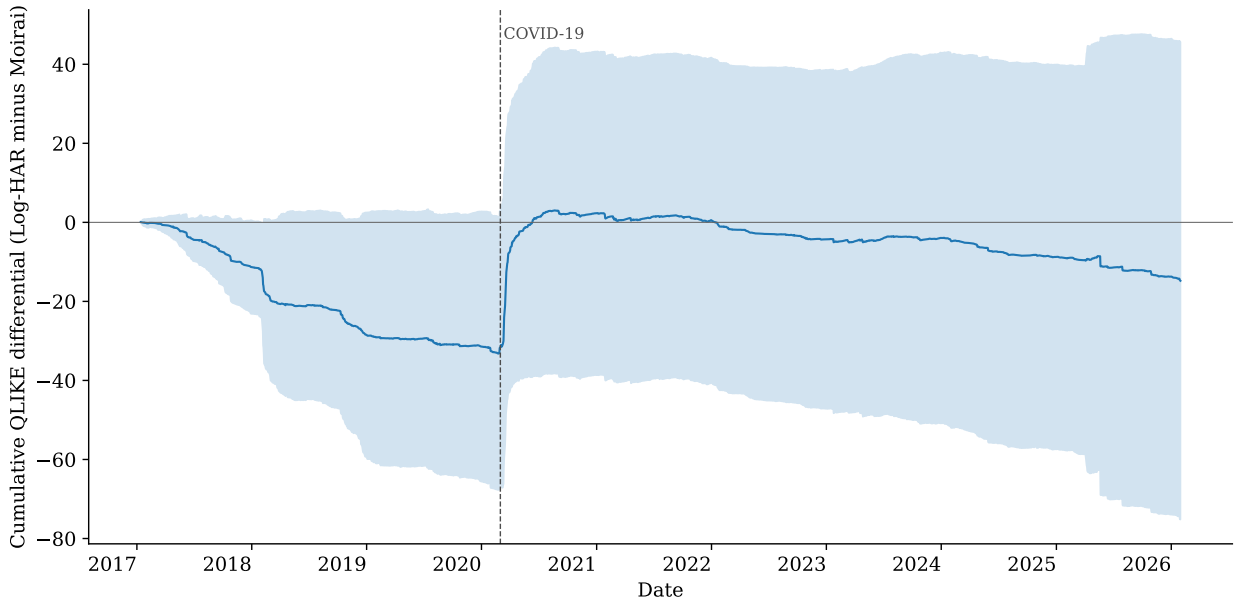

Figure 4: Cumulative QLIKE loss differential (Log-HAR minus Moirai 2.0) at $h = 1$, averaged across 40 equities. Positive values indicate that Moirai 2.0 outperforms Log-HAR. The shaded area is ± 1 cross-sectional standard deviation. The vertical dashed line marks March 2020 (COVID-19 onset).

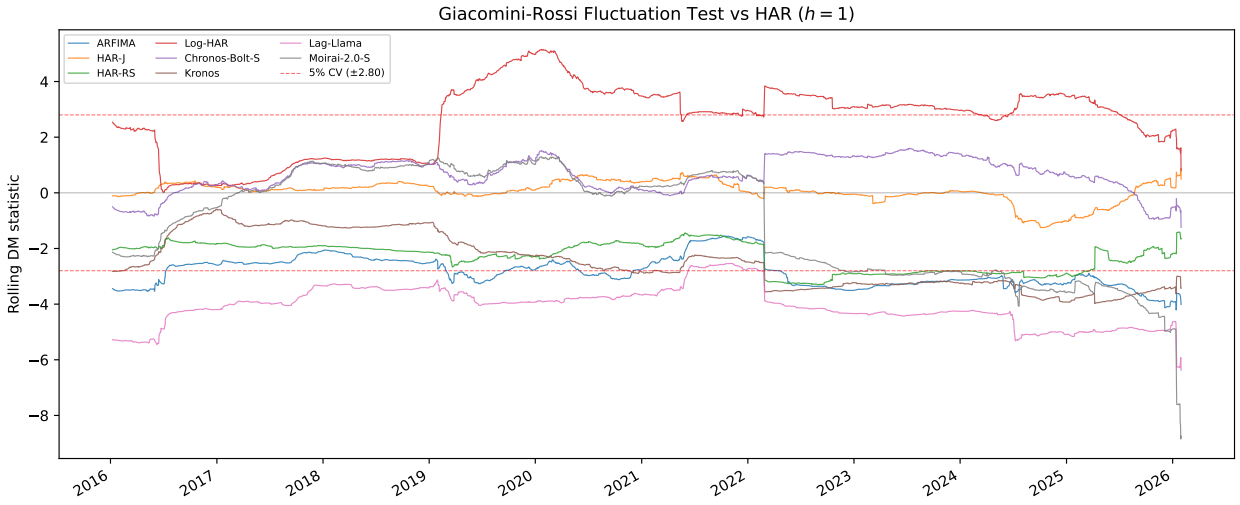

Figure 5: Giacomini–Rossi fluctuation test: rolling DM statistic (QLIKE loss) of each model vs. HAR at $h = 1$, averaged across 50 assets. Negative values indicate the comparison model outperforms HAR. Dashed red lines mark the 5% critical values (± 2.80). The window size is 30% of the sample.

Asset	HAR	Log-HAR	ARFIMA	Chr-Bolt-S	Chr-Bolt-B	Moirai-S	Lag-Llama	Kronos
Stocks								
AAPL	0.1997	**0.1932**	0.4680	0.2273	0.2321	0.2061	0.4127	0.3365
ADBE	0.1570	**0.1507**	0.4303	0.1755	0.1789	0.1617	0.3462	0.4312
AMD	0.1532	**0.1500**	0.2867	0.1692	0.1720	0.1630	0.2877	0.3163
AMGN	**0.2145**	0.2155	0.3355	0.2422	0.2369	0.2210	0.3878	0.9990
AMZN	12.5968	**0.4498**	0.7394	0.5050	0.5127	0.4790	0.8656	10.9395
AXP	0.2102	0.2168	0.3013	0.2111	0.2106	**0.1959**	0.4140	0.3952
BA	8.4780	0.2750	0.3601	0.2439	0.2447	**0.2248**	0.4943	0.8086
C	0.3190	**0.2943**	0.3345	0.3256	0.3269	0.2969	0.5447	1.9342
CAT	0.1755	0.1753	0.2539	0.1895	0.1907	**0.1693**	0.3422	0.2478
CRM	0.2429	**0.2208**	0.3825	0.2516	0.2520	0.2341	0.3983	1.1444
CSCO	0.1910	**0.1829**	0.3138	0.2102	0.2098	0.1913	0.3785	0.3046
CVX	0.1810	0.1867	0.2713	0.1583	**0.1566**	0.1658	0.3065	0.2955
DIS	2.3357	0.2159	0.2996	0.2302	0.2345	**0.2058**	0.4009	0.3025
GE	0.2165	**0.1918**	0.2825	0.2173	0.2185	0.2025	0.3640	0.8605
GOOGL	0.3823	0.3179	0.4568	0.3416	0.3519	**0.3140**	0.5129	0.9987
GS	0.1639	0.1596	0.2705	0.1739	0.1715	**0.1529**	0.3189	0.2449
HD	0.2076	0.1944	0.3010	0.1957	0.1939	**0.1836**	0.4202	0.4611
HON	0.2257	0.2181	0.3087	0.2078	**0.2058**	0.2115	0.4032	0.3393
IBM	0.6349	**0.1923**	0.3080	0.2133	0.2157	0.2109	0.3832	0.2483
JNJ	0.2686	0.2236	0.2970	0.2152	**0.2128**	0.2529	0.3978	1.9198
JPM	0.1908	0.1818	0.3034	0.1950	0.1943	**0.1768**	0.3618	0.3269
KO	0.1855	0.1732	0.2471	0.1675	**0.1655**	0.2524	0.3446	0.2906
MCD	3.5731	0.2319	0.2942	**0.2252**	0.2259	0.2622	0.3831	0.4779
META	0.2420	0.1929	0.3438	0.1887	0.1951	**0.1790**	0.3721	0.5698
MMM	0.2183	**0.2062**	0.2949	0.2271	0.2280	0.2509	0.3729	0.3599
MRK	0.2077	0.2008	0.2712	0.2212	0.2184	**0.1922**	0.3507	0.2930
MSFT	0.2917	**0.2616**	0.4086	0.3311	0.3256	0.3193	0.4759	1.1445
NFLX	9.5971	**0.3634**	0.4565	0.3809	0.3801	0.3696	0.5480	7.3007
NKE	0.2210	0.2053	0.2797	0.2105	0.2125	**0.1818**	0.3396	0.5757
NVDA	0.1891	**0.1813**	0.3218	0.2017	0.2062	0.1907	0.3857	0.6333
ORCL	1.5189	**0.2516**	0.3804	0.2854	0.2961	0.2608	0.5327	0.8466
PG	0.2115	0.1940	0.2596	0.1812	**0.1806**	0.2332	0.3771	0.4822
PM	0.2546	**0.2384**	0.2913	0.2549	0.2589	0.2433	0.4061	0.3556
SHW	1.3864	0.2082	0.2602	0.1991	0.2005	**0.1784**	0.3264	0.2478
TRV	0.1709	**0.1655**	0.2357	0.1726	0.1691	0.1858	0.3279	0.2118
TSLA	2.0318	**0.1716**	0.2922	0.2029	0.2036	0.1869	0.3343	0.3639
UNH	0.2342	**0.2088**	0.3070	0.2519	0.2527	0.2228	0.4056	0.4357
V	0.7542	**0.2523**	0.3078	0.2686	0.2732	0.3309	0.4384	3.9474
VZ	1.0678	**0.1814**	0.2630	0.1998	0.2035	0.2040	0.3847	0.3659
WMT	0.8794	**0.2470**	0.3088	0.2517	0.2498	0.2743	0.4637	2.5587
XOM	0.1999	0.2174	0.2622	0.1709	**0.1678**	0.2122	0.3034	0.3291
Average	1.2483	**0.2185**	0.3266	0.2315	0.2326	0.2281	0.4052	1.1133
FX								
AUDUSD	0.1572	**0.1483**	0.2116	0.1504	0.1536	0.3086	0.2664	0.3134
EURUSD	0.1516	**0.1441**	0.1980	0.1576	0.1594	0.3139	0.2768	0.2266
GBPUSD	0.1451	**0.1334**	0.2035	0.1492	0.1487	0.3198	0.2806	0.2278
USDCAD	0.1351	**0.1282**	0.1725	0.1326	0.1331	0.2587	0.2376	0.1903
USDJPY	0.2720	**0.2415**	0.3957	0.2823	0.2856	0.4456	0.4621	0.5393
Average	0.1722	**0.1591**	0.2363	0.1744	0.1761	0.3293	0.3047	0.2995
Futures								
CL	0.1950	**0.1351**	0.3268	0.1591	0.1521	0.1741	3.0510	45.4366
ES	0.2234	**0.1928**	0.5353	0.2244	0.2224	0.5307	1.2473	0.7401
GC	0.2574	**0.1626**	0.2860	0.1763	0.2223	0.2669	0.3741	1.2605
NG	0.6772	0.1800	0.2402	0.1790	0.1802	**0.1773**	0.3715	3.4776
Average	0.3383	**0.1676**	0.3471	0.1847	0.1942	0.2873	1.2610	12.7287
Overall Average	1.0679	**0.2085**	0.3192	0.2221	0.2239	0.2429	0.4636	1.9611

Per-asset QLIKE loss at $h = 1$. QLIKE is computed as $\text{mean}(RV_t/\hat{RV}_t - \ln(RV_t/\hat{RV}_t) - 1)$ with forecasts floored at 10^{-6}. Bold indicates the lowest (best) QLIKE in each row. Chr-Bolt-S = Chronos-Bolt-Small, Chr-Bolt-B = Chronos-Bolt-Base, Moirai-S = Moirai-2.0-Small. HAR-J, HAR-RS, and HARQ are omitted as their QLIKE values are dominated by near-zero forecast artifacts (see Section 5.1).

Model	$\hat{\alpha}$	$\hat{\beta}$	R^2	F-stat	p(joint)	% Reject
Econometric benchmarks						
HAR	0.0000	1.155	0.276	7.09	0.261	28%
HAR-J	0.0000	1.154	0.293	5.99	0.268	28%
HAR-RS	0.0001	0.910	0.210	15.64	0.143	52%
HARQ	0.0001	0.681	0.119	13.01	0.063	70%
Log-HAR	−0.0000	1.438	0.198	5.42	0.050	76%
ARFIMA	0.0001	0.977	0.221	15.90	0.099	72%
Foundation models						
Chronos-Bolt-S	0.0000	1.282	0.400	26.29	0.002	98%
Chronos-Bolt-B	0.0000	1.438	0.417	31.24	0.001	98%
Moirai-2.0-S	0.0000	1.094	0.367	18.12	0.007	94%
Lag-Llama	0.0000	1.183	0.069	5.01	0.055	60%
Kronos	0.0000	1.298	0.295	142.51	0.000	100%

Mincer–Zarnowitz forecast efficiency regression, $h = 5$. See Table 5.3 for variable definitions.

Model	$\hat{\alpha}$	$\hat{\beta}$	R^2	F-stat	p(joint)	% Reject
Econometric benchmarks						
HAR	0.0001	0.490	0.096	18.82	0.044	70%
HAR-J	0.0001	0.510	0.102	16.83	0.044	70%
HAR-RS	0.0002	0.366	0.083	32.16	0.016	92%
HARQ	0.0002	0.299	0.074	54.19	0.003	100%
Log-HAR	0.0001	0.474	0.072	16.70	0.018	92%
ARFIMA	0.0001	0.578	0.143	16.87	0.040	80%
Foundation models						
Chronos-Bolt-S	0.0001	1.121	0.202	10.98	0.008	96%
Chronos-Bolt-B	0.0000	1.453	0.221	12.15	0.005	98%
Moirai-2.0-S	0.0000	0.985	0.187	12.55	0.006	98%
Lag-Llama	0.0001	0.705	0.057	3.43	0.096	50%
Kronos	0.0001	0.906	0.127	24.00	0.001	100%

Mincer–Zarnowitz forecast efficiency regression, $h = 22$. See Table 5.3 for variable definitions.

B Supplementary Figures

Figures 6 and 7 extend the Giacomini–Rossi fluctuation test analysis (Figure 5) to the $h = 5$ and $h = 22$ horizons.

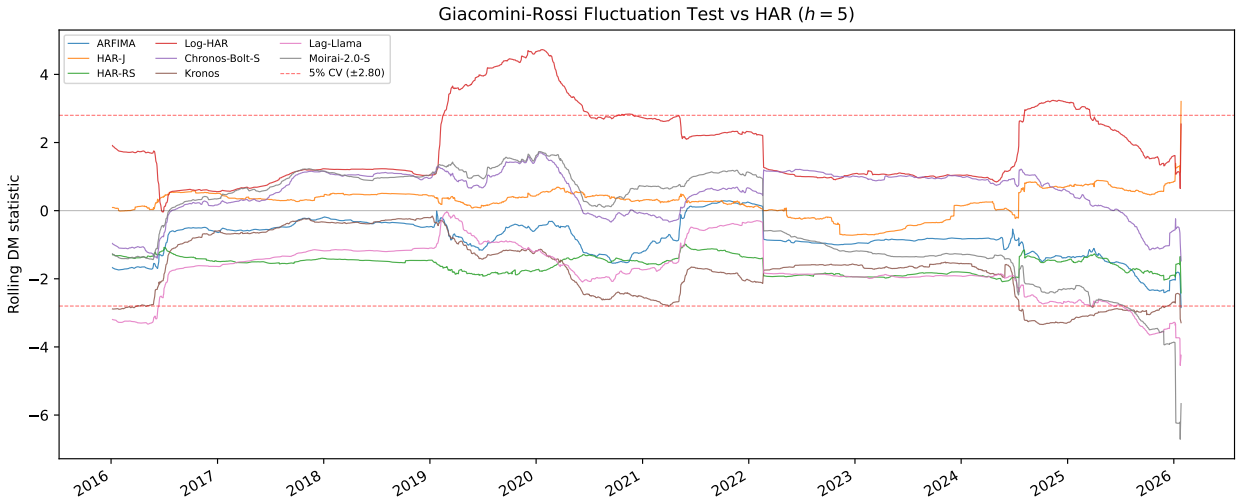

Figure 6: Giacomini–Rossi fluctuation test at $h = 5$. Each panel plots the rescaled cumulative loss differential between the column model and HAR, with $\pm 5\%$ critical values (dashed).

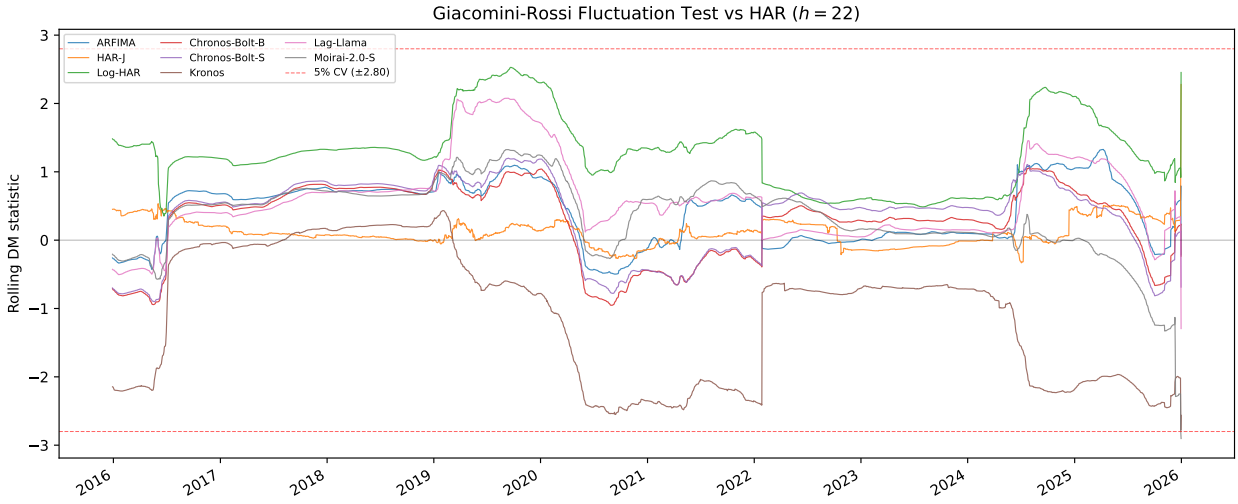

Figure 7: Giacomini–Rossi fluctuation test at $h = 22$. Each panel plots the rescaled cumulative loss differential between the column model and HAR, with $\pm 5\%$ critical values (dashed).